



Text Processing using Machine Learning

Introduction: Classic vs Deep NLP

Liling Tan

02 Dec 2019

OVER
5,500 GRADUATE
ALUMNI

OFFERING OVER
120 ENTERPRISE IT, INNOVATION
& LEADERSHIP PROGRAMMES

TRAINING OVER
120,000 DIGITAL LEADERS
& PROFESSIONALS

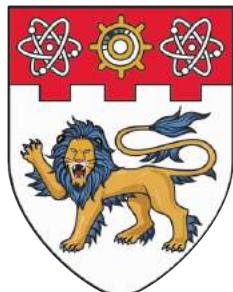
About Me

Work:

Rakuten
Institute of Technology

<https://rit.rakuten.co.jp/about/>

Education:



NANYANG
TECHNOLOGICAL
UNIVERSITY
SINGAPORE



UNIVERSITÄT
DES
SAARLANDES

User:Alvations

From Wikipedia, the free encyclopedia

食飽未?

Alvas (short for alvations) is a geek with a hint of linguistics, dabbling in philosophy. Occam's razor is his rule of thumb and sometimes he wanders into his Truman's land or his Norwegian Wood.

Currently, alvas is working on machine translation, language learning technologies and Natural Language Processing at Rakuten Institute of Technology. Previously, alvas was working on translation technologies(human and machine) in the Translation Technology Group at the Sprachwissenschaften sowie Übersetzen und Dolmetschen department @ Universität des Saarlandes. Previous previously, alvas was a linguistics graduate student in Bond lab @ NTU (ミニオンは今までたってもミニオン). In his linguistics pursuit, he would love to work with his mother tongue, Hokkien. Meanwhile he's trying to let computers understand and produce human language and spreading Awesomeness.

Currently, alvas seldom edit stuff on Wikipedia, but here's some things that alvas did on wikipedia:

- Setup a wikipedia page for SemEval during the "Language and Computer" course

This user is a [linguist](#).
...!
re-0
prog-N
This user activates the [Stargate](#) to invade other planets.
This user is not afraid of the [Evil Monkey](#).
This user knows it all started with a [Big Bang \(BANG!!!\)](#)

What is this course about?

- **Recent advance techniques for NLP**
 - esp. since the “deep learning tsunami”
- Getting a **good grasp of PyTorch for NLP**
 - Similar knowledge can be applied to any tensor libraries for machine/deep learning
- **Understanding the underlying techniques and knowing how to implement them**
 - Hopefully, you’ll find them less “magical”

What is this course NOT about?

- **Understanding everything in NLP.**
 - There's a lot more to computational linguistics than what's in this course
- Using **libraries to achieve results for specific NLP tasks**
 - (e.g. SpaCy, Transformer, Gluon, AllenNLP, TorchText, etc.)
- No, **we won't build chatbots** in this course.

Pardon my quirks...

- **Sometimes I let the code do the talking and I just read out the code** (*Stop me if it's not intuitive enough*)
- **I have more code than math in my slides**
- **Feel free to stop me at any time to ask questions**

Some Questions...

- How many people learnt deep / machine learning from previous courses (in-class or MOOC)?
- How many of from computer science/engineering background? How many from humanities, STEM, non-computer engineering or business background? Others?
- How many currently working in a technology field? Industry/Academia?

Course Schedule

Introduction

Classic vs Deep NLP

NN from Scratch

Deep Learning Foundations

Matrix Calculus for Deep Learning

Backpropagation

Word Embeddings

Word2Vec, GloVe, Fasttext, and friends

Word Embeddings from Scratch

Nuts and Bolts

Bias - Variance

Regularization, Loss Functions, Optimizers

Language Models

N-gram + Neural Language Models

Recurrent Neural Nets

Memory Networks and Conditional Generation

Exploding and Vanishing Gradients

LSTM + GRU and Seq2Seq Models

Attention Networks

Lots of Different Attentions

Attention is all you need (aka "Transformer")

Sentence Representation

Multi-tasks vs Transfer Learning

Pretrained Language Models (aka "Sesame Street")

Machine Translation (Bonus)

Phrase-based to Neural MT

101 tricks to Neural MT training and quirks

Summary

Recap

Ask Me Anything

Course Logistics

For the hands-on later, it'll take a while to download and install, so we do this first while the lesson goes on...

Go to <https://www.anaconda.com/download>

Download the Python3 version and install

Overview

Lecture

- Classic NLP
- Deep Magic NLP
- Deep Learning Basics

Hands-on

- Environment Setup
- Deep Learning From Scratch

Classic NLP

Vector space model, Frequency, TF-IDF and PPMI



You shall know a word by the company it keeps...

– John R. Firth (1957)



"You shall know a word by the company it keeps..."
– John R. Firth (1957)

***"Everyone is about Firth 1957 (you shall know a word...).
Somehow we all skipped Firth 1935"***
– Yoav Goldberg (2016)

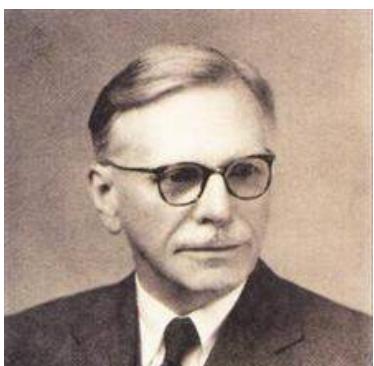




"You shall know a word by the company it keeps..."
– John R. Firth (1957)



***"Everyone is about Firth 1957 (you shall know a word...).
Somehow we all skipped Firth 1935"***
– Yoav Goldberg (2016)



***"... the complete meaning of a word is always
contextual, and no study of meaning apart from
context can be taken seriously"***
– John R. Firth (1935)



“The frequencies of word in a document tend to indicate the relevance of document to a query”
– Gerard Salton (1975)

“From frequency to meaning.... Statistical patterns of human word usage can be used to figure out what people mean”
– Turney and Pantel (2010)



Vector Space Model

- Vector space models are **numerical representation of text**
- Traditionally, computed using **no. of times each word occurs**

Frequency (Count Vector)

```
sent0 = "The quick brown fox jumps over the lazy brown dog ."
```

```
sent1 = "Mr brown jumps over the lazy fox ."
```

	brown	dog	fox	jumps	lazy	mr	over	quick	the
sent0	2	1	1	1	1	0	1	1	2
sent1	1	0	1	1	1	1	1	0	1

Term Frequency – Inverse Document Frequency

$$tf_{t,d} = \begin{cases} 1 + \log_{10} \text{count}(t,d) & \text{if } \text{count}(t,d) > 0 \\ 0 & \text{otherwise} \end{cases}$$

$$idf_i = \log_{10} \frac{N}{df_i}$$

total no. of sentences

no. of sentences that contains word i

tf-idf value for word t in document d :

$$w_{t,d} = tf_{t,d} \times idf_t$$

```
1 from sklearn.feature_extraction.text import TfidfVectorizer
2 import numpy as np
3 from scipy.sparse.csr import csr_matrix
4
5 # The *TfidfVectorizer* from sklearn expects list of strings as input.
6 sent0 = "The quick brown fox jumps over the lazy brown dog .".lower()
7 sent1 = "Mr brown jumps over the lazy fox .".lower()
8
9 dataset = [sent0, sent1]
10
11 vectorizer = TfidfVectorizer(input=dataset, analyzer='word',
12                             ngram_range=(1,1), min_df = 0, stop_words=None)
13 tfidf_matrix = vectorizer.fit_transform(dataset)
14
15 # Format the TF-IDF table into the pd.DataFrame format.
16 vocab = vectorizer.get_feature_names()
17 documents_tfidf_lol = [{word:tfidf_value
18                         for word, tfidf_value in zip(vocab, sent)}
19                         for sent in tfidf_matrix.toarray()]
20 documents_tfidf = pd.DataFrame(documents_tfidf_lol)
21 documents_tfidf.fillna(0, inplace=True)
```

Term Frequency – Inverse Document Frequency

```
sent0 = "The quick brown fox jumps over the lazy brown dog ."
```

```
sent1 = "Mr brown jumps over the lazy fox ."
```

	brown	dog	fox	jumps	lazy	mr	over	quick	the
sent0	0.500	0.351	0.250	0.250	0.250	0.000	0.250	0.351	0.500
sent1	0.354	0.000	0.354	0.354	0.354	0.497	0.354	0.000	0.354

Classification (Train/Test Split)

Train

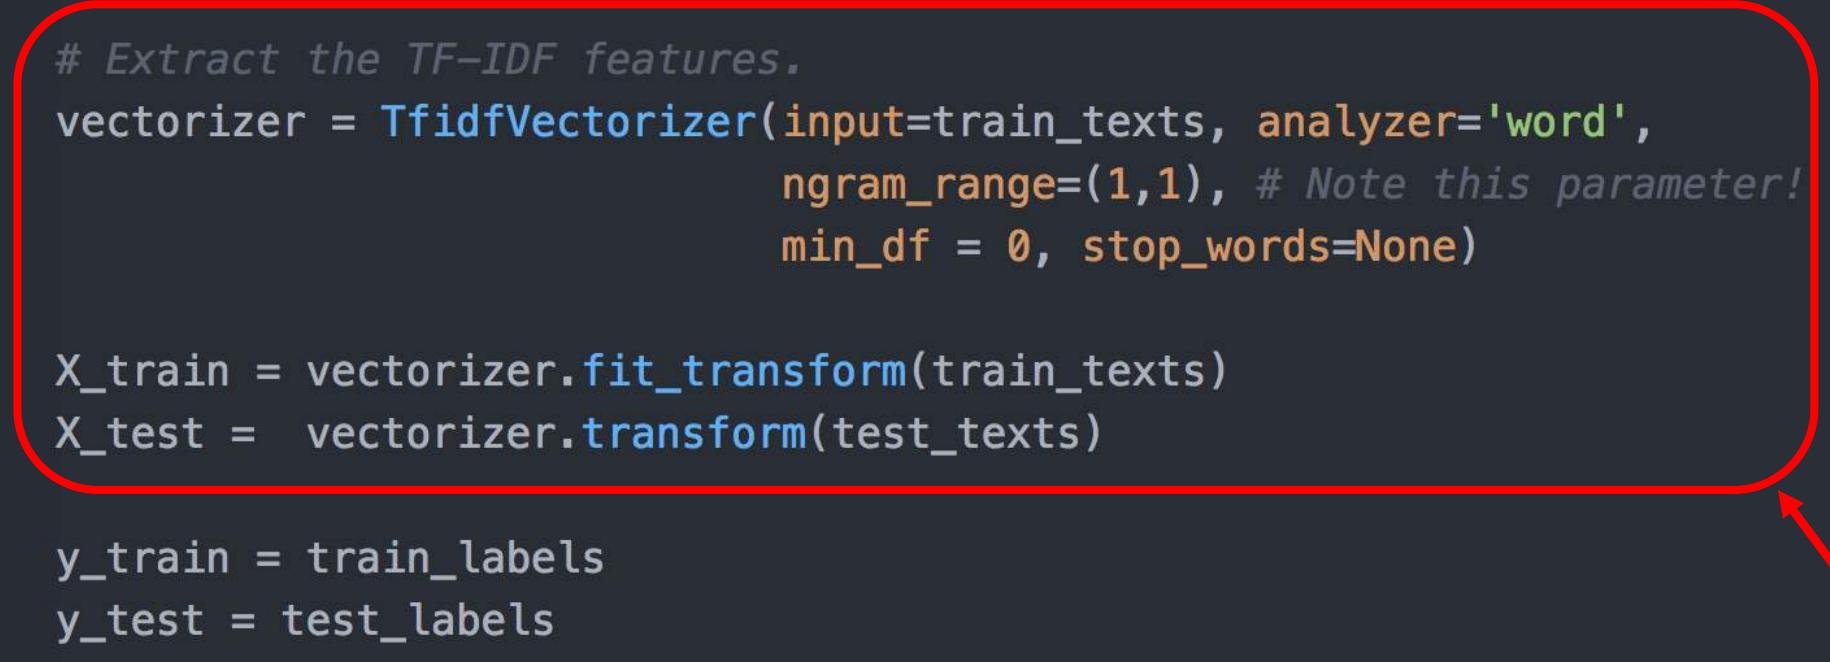
```
sent0 = "The quick brown fox jumps over the lazy brown dog ."  
sent1 = "Mr brown jumps over the lazy fox ."  
sent2 = "Roses are red , the chocolates are brown ."  
sent3 = "The frank dog jumps through the red roses ."
```

Test

```
sent4 = "Mr Tan jumps on red chocolates ?"  
sent5 = "Mr brown likes the lazy dog ."
```

```
24 # Train sentences.  
25 sent0, label0 = "The quick brown fox jumps over the lazy brown dog .".lower() , False  
26 sent1, label1 = "Mr brown jumps over the lazy fox .".lower(), True  
27 sent2, label2 = "Roses are red , the chocolates are brown .".lower(), False  
28 sent3, label3 = "The frank dog jumps through the red roses .".lower(), False  
29  
30 # Test sentences.  
31 sent4, label4 = "Mr Tan jumps on red chocolates ?".lower(), False  
32 sent5, label5 = "Mr brown likes the lazy dog .".lower(), True  
33  
34 train_documents = [(sent0, label0), (sent1, label1),  
35 (sent2, label2), (sent0, label2)]  
36 test_documents = [(sent4, label4), (sent5, label5)]  
37  
38 train_texts, train_labels = zip(*train_documents)  
39 test_texts, test_labels = zip(*test_documents)
```

```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the TF-IDF features.
48 vectorizer = TfidfVectorizer(input=train_texts, analyzer='word',
49                             ngram_range=(1,1), # Note this parameter!
50                             min_df = 0, stop_words=None)
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```



Apply TF-IDF
to the train
and test set

```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the TF-IDF features.
48 vectorizer = TfidfVectorizer(input=train_texts, analyzer='word',
49                               ngram_range=(1,1), # Note this parameter!
50                               min_df = 0, stop_words=None)
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```

Apply the
Machine
Learning
algorithm



```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the TF-IDF features.
48 vectorizer = TfidfVectorizer(input=train_texts, analyzer='word',
49                               ngram_range=(1,1), # Note this parameter!
50                               min_df = 0, stop_words=None)
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```

Accuracy
= 50%



```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the TF-IDF features.
48 vectorizer = TfidfVectorizer(input=train_texts, analyzer='word',
49                             ngram_range=(2,2), # Note this parameter!
50                             min_df = 0, stop_words=None)
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```

Lets use TF-IDF with **bigrams**

Term Frequency – Inverse Document Frequency

```
sent0 = "The quick brown fox jumps over the lazy brown dog ."
```

```
sent1 = "Mr brown jumps over the lazy fox ."
```

	brown dog	brown fox	brown jumps	fox jumps	jumps over	...	mr brown	the lazy	the quick
sent0	0.364	0.364	0.000	0.364	0.259	...	0.000	0.259	0.364
sent1	0.000	0.000	0.470	0.000	0.334	...	0.470	0.334	0.000

```

43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the TF-IDF features.
48 vectorizer = TfidfVectorizer(input=train_texts, analyzer='word',
49                             ngram_range=(2,2), # Note this parameter!
50                             min_df = 0, stop_words=None)
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))

```

Accuracy
 = 100%

But we sort of
 “cheated” and
 we engineered
 the TD-IDF
 features

Pointwise Mutual Information

Does the w_1 and w_2 co-occur more than if they were independent?

$$\text{PMI}(w_1, w_2) = \log_2 \frac{P(w_1, w_2)}{P(w_1)P(w_2)}$$

Pointwise Mutual Information

Does the w_1 and w_2 co-occur more than if they were independent?

$$\text{PMI}(w_1, w_2) = \log_2 \frac{P(w_1, w_2)}{P(w_1)P(w_2)}$$

$$P_{*j} = \frac{\sum_{i=1}^W f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}} \quad P_{i*} = \frac{\sum_{j=1}^C f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}} \quad P_{ij} = \frac{f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}} \quad \text{PMI}_{ij} = \log_2 \frac{p_{ij}}{p_{i*} p_{*j}}$$

Co-occurrence Matrix

```
sent1 = "Mr brown jumps over mr fox ."
```

	mr	brown	jumps	over	fox
mr					
brown					
jumps					
over					
fox					

Co-occurrence Matrix

```
sent1 = "Mr brown jumps over mr fox ."
```

	mr	brown	jumps	over	fox
mr		1	1		
brown					
jumps					
over					
fox					

Co-occurrence Matrix

```
sent1 = "Mr brown jumps over mr fox ."
```

	mr	brown	jumps	over	fox
mr		1	1		
brown			1	1	
jumps					
over					
fox					

Co-occurrence Matrix

```
sent1 = "Mr brown jumps over mr fox ."
```

	mr	brown	jumps	over	fox
mr		1	1		
brown			1	1	
jumps	1			1	
over					
fox					

Co-occurrence Matrix

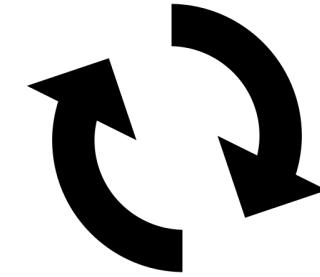
```
sent1 = "Mr brown jumps over mr fox ."
```

	mr	brown	jumps	over	fox
mr		1	1		
brown				1	1
jumps	1				1
over	1				1
fox					

Co-occurrence Matrix

```
sent2 = "Fox jumps over mr brown ."
```

	mr	brown	jumps	over	fox
mr		1	1		
brown			1	1	
jumps	1			1	
over	1				1
fox					



Repeat the same process for all sentences and add to the counts in the cells

Pointwise Mutual Information

		$j = 1$	$j = 2$	$j = 3$	$j = 4$	$j = 5$
		mr	brown	jumps	over	fox
$i = 1$	mr	0	1	1	0	0
$i = 2$	brown	0	0	1	1	2
$i = 3$	jumps	1	0	0	2	1
$i = 4$	over	2	1	0	0	1
$i = 5$	fox	0	0	1	0	0

Matrix F with

- W rows (words)

- C columns (context)

Pointwise Mutual Information

		$j = 1$	$j = 2$	$j = 3$	$j = 4$	$j = 5$
		mr	brown	jumps	over	fox
$i = 1$	mr	0	1	1	0	0
$i = 2$	brown	0	0	1	1	2
$i = 3$	jumps	1	0	0	2	1
$i = 4$	over	2	1	0	0	1
$i = 5$	fox	0	0	1	0	0

Matrix F with
 - W rows (words)
 - C columns (context)

$$\sum_{i=1}^W \sum_{j=1}^C f_{ij} = 15$$

Pointwise Mutual Information

		$j = 1$	$j = 2$	$j = 3$	$j = 4$	$j = 5$
		mr	brown	jumps	over	fox
$i = 1$	mr	0	1	1	0	0
$i = 2$	brown	0	0	1	1	2
$i = 3$	jumps	1	0	0	2	1
$i = 4$	over	2	1	0	0	1
$i = 5$	fox	0	0	1	0	0

$$\sum_{i=1}^W \sum_{j=1}^C f_{ij} = 15$$

Matrix F with
 - W rows (words)
 - C columns (context)

$$P_{ij} = \frac{f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P(jumps, over) = 2 / 15$$

Pointwise Mutual Information

		$j = 1$	$j = 2$	$j = 3$	$j = 4$	$j = 5$
		mr	brown	jumps	over	fox
$i = 1$	mr	0	1	1	0	0
$i = 2$	brown	0	0	1	1	2
$i = 3$	jumps	1	0	0	2	1
$i = 4$	over	2	1	0	0	1
$i = 5$	fox	0	0	1	0	0

$$\sum_{i=1}^W \sum_{j=1}^C f_{ij} = 15$$

$$\sum_{j=1}^C f_{ij} = 4$$

Matrix F with
 - W rows (words)
 - C columns (context)

$$P_{i^*} = \frac{\sum_{j=1}^C f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P(jumps, over) = 2 / 15$$

$$P(jumps) = 4 / 15$$

Pointwise Mutual Information

		$j = 1$	$j = 2$	$j = 3$	$j = 4$	$j = 5$
		mr	brown	jumps	over	fox
$i = 1$	mr	0	1	1	0	0
$i = 2$	brown	0	0	1	1	2
$i = 3$	jumps	1	0	0	2	1
$i = 4$	over	2	1	0	0	1
$i = 5$	fox	0	0	1	0	0

$$\sum_{i=1}^W \sum_{j=1}^C f_{ij} = 15 \quad \sum_{j=1}^C f_{ij} = 4 \quad \sum_{i=1}^W f_{ij} = 3$$

Matrix F with
 - W rows (words)
 - C columns (context)

$$P_{*j} = \frac{\sum_{i=1}^W f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P(jumps, over) = 2 / 15 \\ P(jumps) = 4 / 15 \\ P(over) = 3 / 15$$

Pointwise Mutual Information

	mr	brown	jumps	over	fox
mr	0	1	1	0	0
brown	0	0	1	1	2
jumps	1	0	0	1.321	1
over	2	1	0	0	1
fox	0	0	1	0	0

$$P(jumps, over) = 2 / 15$$

$$P(jumps) = 4 / 15$$

$$P(over) = 3 / 15$$

$$PMI(jumps, over) = \log_2 \frac{P(jumps, over)}{P(jumps) * P(over)} = 1.321$$

$$P_{*j} = \frac{\sum_{i=1}^W f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P_{i*} = \frac{\sum_{j=1}^C f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P_{ij} = \frac{f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$PMI_{ij} = \log_2 \frac{P_{ij}}{P_{i*} P_{*j}}$$

Positive Pointwise Mutual Information

	mr	brown	jumps	over	fox
mr	0	1	1	0	0
brown	0	0	1	1	2
jumps	1	0	0	1.321	1
over	2	1	0	0	1
fox	0	0	1	0	0

$$P(jumps, over) = 2 / 15$$

$$P(jumps) = 4 / 15$$

$$P(over) = 3 / 15$$

$$\text{PMI}(jumps, over) = 1.321$$

$$P_{*j} = \frac{\sum_{i=1}^W f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P_{i*} = \frac{\sum_{j=1}^C f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$P_{ij} = \frac{f_{ij}}{\sum_{i=1}^W \sum_{j=1}^C f_{ij}}$$

$$PMI_{ij} = \log_2 \frac{p_{ij}}{p_{i*} p_{*j}}$$

$$PPMI_{ij} = \begin{cases} PMI_{ij} & \text{if } PMI_{ij} > 0 \\ 0 & \text{Otherwise} \end{cases}$$

Pointwise Mutual Information

```
>>> import math  
  
>>> math.log2(50/(100*200))  
-8.643856189774725
```

$$PMI_{ij} = \log_2 \frac{p_{ij}}{p_i * p_{*j}}$$

Pointwise Mutual Information

```
>>> import math  
  
>>> math.log2(50/(100*200))  
-8.643856189774725  
  
>>> math.log2(5/(10*20))  
-5.321928094887363
```

$$PMI_{ij} = \log_2 \frac{p_{ij}}{p_i * p_{*j}}$$

(higher is better)



If we simply plug-in
words with similar ratio

Pointwise Mutual Information

```
>>> import math  
  
>>> math.log2(50/(100*200))  
-8.643856189774725  
  
>>> math.log2(5/(10*20))  
-5.321928094887363  
  
>>> math.log2(1/(2*4))  
-3.0
```

$$PMI_{ij} = \log_2 \frac{p_{ij}}{p_i * p_{*j}}$$

(higher is better)

*With the same ratio,
PMI is biased towards
rare words!!*

```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the TF-IDF features.
48 vectorizer = TfidfVectorizer(input=train_texts, analyzer='word',
49                             ngram_range=(2,2), # Note this parameter!
50                             min_df = 0, stop_words=None)
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```

```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the PPMI features.
48 vectorizer =PPMIVectorizer(input=train_texts, analyzer='word',
49                             context_window=2, stop_words=None)
50
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```

Replace the
TF-IDF
features with
PPMI*

* There's no
PPMIVectorizer in
sklearn, you have to
write it yourself

```
43 from sklearn.feature_extraction.text import TfidfVectorizer
44 import numpy as np
45 from scipy.sparse.csr import csr_matrix
46
47 # Extract the PPMI features.
48 vectorizer =PPMIVectorizer(input=train_texts, analyzer='word',
49                             context_window=2, stop_words=None)
50
51
52 X_train = vectorizer.fit_transform(train_texts)
53 X_test = vectorizer.transform(test_texts)
54
55 y_train = train_labels
56 y_test = test_labels
57
58 # Pick your poison.
59 from sklearn.linear_model import Perceptron
60 # Initialize your classifier.
61 clf = Perceptron(max_iter=10)
62 # Train the classifier.
63 clf.fit(X_train, y_train)
64
65 print(clf.predict(X_test))
```

Machine
learning block
often don't
change much



- **Raw frequency is useful**
 - but frequent words non-content words are not very informative
- **TF-IDF resolves high freq non-content words issue**
 - but each word is still somewhat *independent* of each other
- **PPMI provides information about whether a word is informative in the context of another word**
 - but biased towards infrequent events

Classic NLP: Feature Engineering

- **TF-IDF and PPMI vectors are**
 - long ($|V| > 100,000$)
 - sparse (lots of zero)
- **Deep learning can create vectors that are**
 - short (often fixed-sized < 2000 , decided empirically)
 - dense (most are non-zeros)
- **But it's not unlike ‘modern’ deep learning based NLP**
 - one model improves upon another often incremental
 - they always come with certain caveats

Deep ‘Magic’ NLP

Recent advancement in Deep Learning in NLP



“The frequencies of word in a document tend to indicate the relevance of document to a query”
– Gerard Salton (1975)

“From frequency to meaning.... Statistical patterns of human word usage can be used to figure out what people mean”
– Turney and Pantel (2010)



Deep ‘Magic’ NLP



“The frequencies of word in a document tend to indicate the relevance of document to a query”
– Gerard Salton (1975)

“From frequency to meaning.... Statistical patterns of human word usage can be used to figure out what people mean”
– Turney and Pantel (2010)



“We propose a unified NN architecture by trying to avoid task-specific engineering therefore disregarding a lot of prior knowledge”
– Collobert and Weston (2011)

Definitions (NLP, ML, DL)

- **Natural Language Processing (NLP) is ...**
 - making computers understand and produce human languages.
- **Machine Learning is ...**
 - optimizing parameters/weights to best make a decision
 - *well-defined counting*
- **Deep Learning, some people say it's ...**
 - neural nets
 - stacking multiple layers of "representation learning"
 - something that burns up as much GPUs as Bitcoin mining
 - a subset of methods in machine learning

The “ImageNet” Moment

Various Computer Vision Challenges (ImageNet, MS Coco, etc.) started a wave of groups **training models and sharing pre-trained models.**

Fine-tuning / transfer learning for these pre-trained models are faster and “usually better” when training new models for other computer vision task.



14,197,122 images, 21841 synsets indexed
[Explore](#) [Download](#) [Challenges](#) [Publications](#) [CoolStuff](#) [About](#)

Not logged in. [Login](#) | [Signup](#)

ImageNet is an image database organized according to the [WordNet](#) hierarchy (currently only the nouns), in which each node of the hierarchy is depicted by hundreds and thousands of images. Currently we have an average of over five hundred images per node. We hope ImageNet will become a useful resource for researchers, educators, students and all of you who share our passion for pictures.
[Click here](#) to learn more about ImageNet, [Click here](#) to join the ImageNet mailing list.



What do these images have in common? *Find out!*

[Check out the ImageNet Challenge on Kaggle!](#)

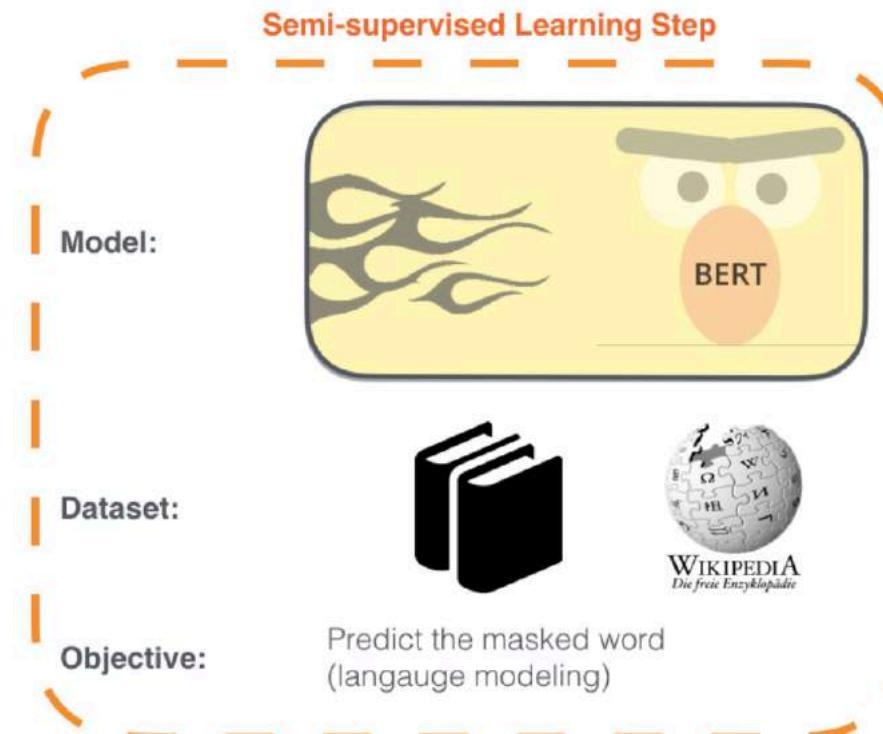
© 2016 Stanford Vision Lab, Stanford University, Princeton University support@image-net.org Copyright infringement

(*Image from [Stanford Vision Lab](#))

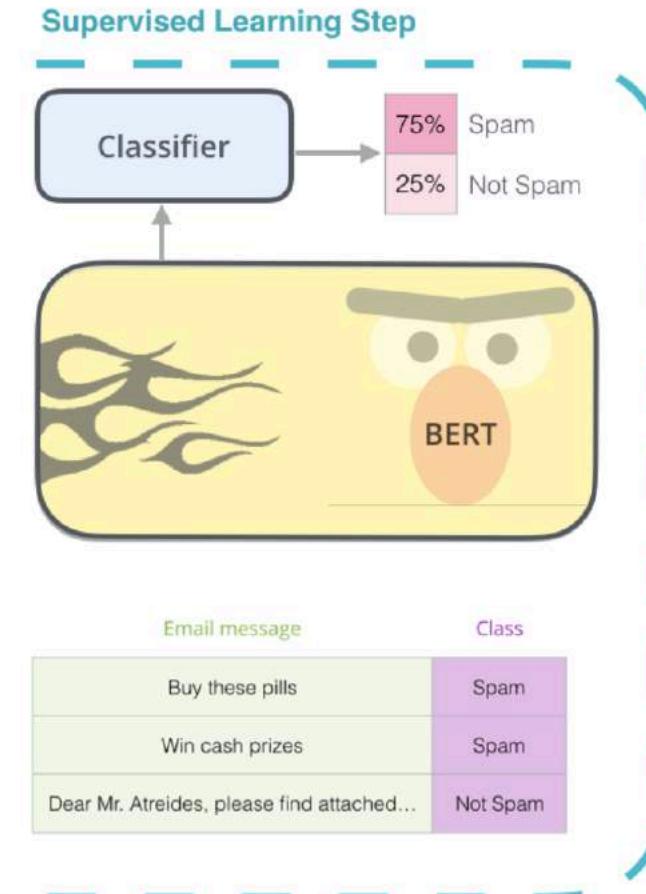
The “ImageNet” Moment for NLP

1 - **Semi-supervised** training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.



2 - **Supervised** training on a specific task with a labeled dataset.



The two steps of how BERT is developed. You can download the model pre-trained in step 1 (trained on un-annotated data), and only worry about fine-tuning it for step 2. [Source for book icon].

(*Image from [Jay Alammar's blog](#))

The “ImageNet” Moment for NLP

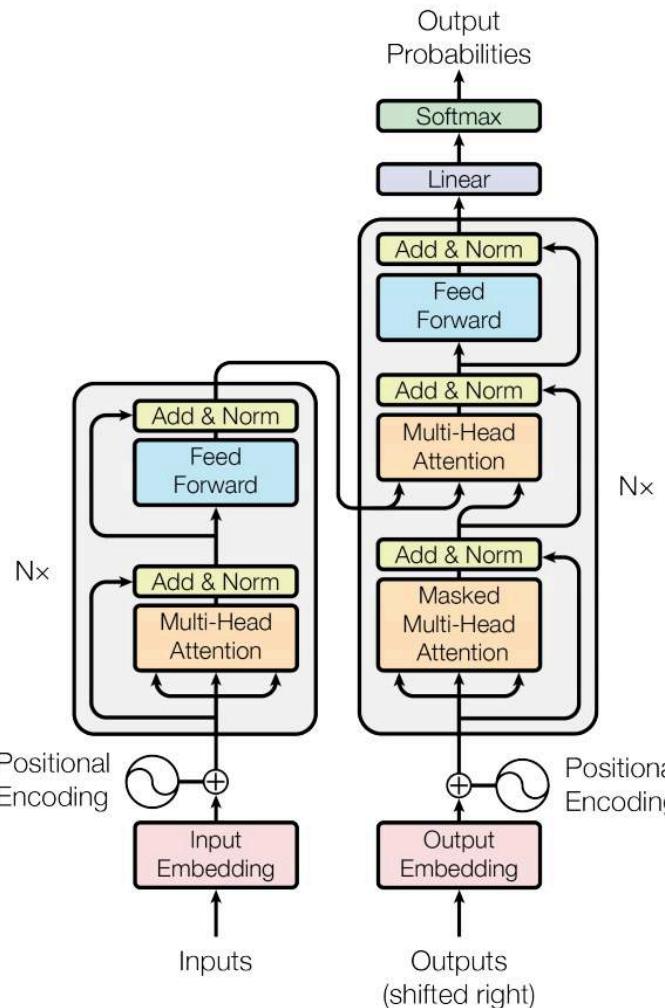


Figure 1: The Transformer - model architecture.

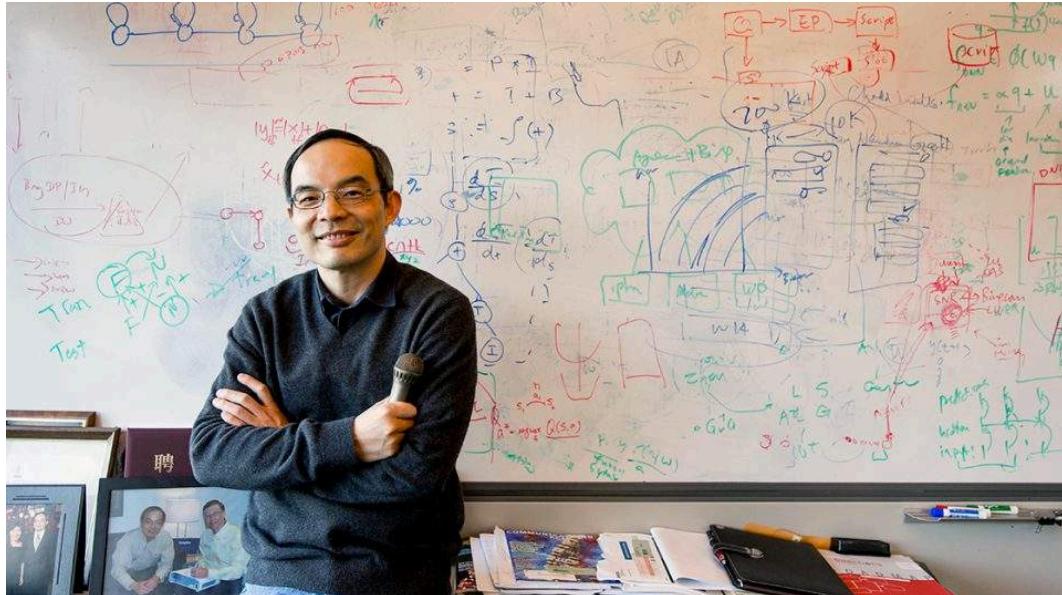
(*Image from Vaswani et al. 2017)

Major breakthrough of BERT came through the **self-attention network architecture, aka. Transformer** (Vaswani et al. 2017)

But do note the caveats with transfer-learning (aka. pre-training and fine-tuning):

- Could get **same results from random initialization** vs pre-training counterparts ([He et al. 2018](#) on “Rethinking ImageNet”)
- Understanding **why pre-training works in NLP still unclear** ([Goldberg, 2018](#), see also [Erhan, 2010](#))

Machine Translation Achieved Human Parity



(*Image from Microsoft AI Blog)

“Microsoft reaches a historic milestone, using AI to match human performance in translating news from Chinese to English”

- [Microsoft AI Blog](#)

Definition of Human Parity

If there is ***no statistically significant difference*** between human quality scores for ... machine translation ... and the scores for the corresponding human translations then the machine has achieved human parity.

Did MT really achieve Human Parity?

- **Has Machine Translation Achieved Human Parity? A Case for Document-level Evaluation**
 - Document level evaluation is necessary
 - <http://aclweb.org/anthology/D18-1512>
- **What Level of Quality can Neural Machine Translation Attain on Literary Text?**
 - Only 17-34% of novels can be machine translated
 - https://link.springer.com/chapter/10.1007/978-3-319-91241-7_12

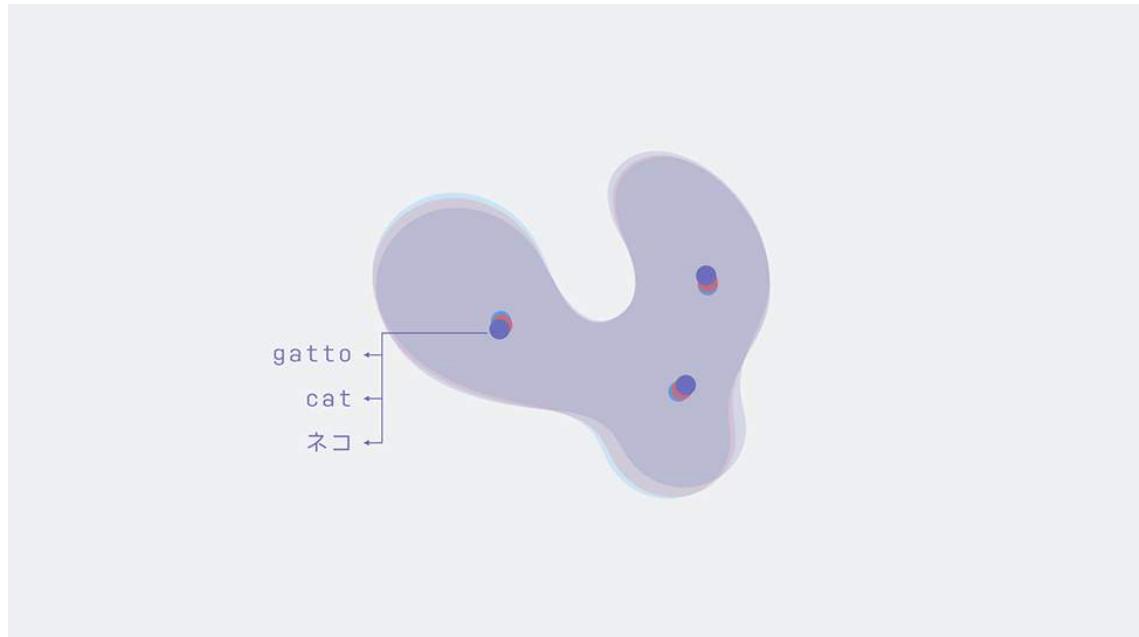
Did MT really achieve Human Parity?

- **Quality expectations of machine translation**
 - “those of us who have seen many paradigms come and go know that overgilding the lily does none of us any good”
 - <https://arxiv.org/pdf/1803.08409.pdf>
- **"Human parity" in machine translation**
 - “Some aspects of such translations are very good, but the frequent mistakes spoil things.”
 - <http://languagelog.ldc.upenn.edu/nll/?p=40602>

Machine Translation Achieved Human Parity

Algorithm 1: Unsupervised MT

```
1 Language models: Learn language models  $P_s$  and  $P_t$   
over source and target languages;  
2 Initial translation models: Leveraging  $P_s$  and  $P_t$ ,  
learn two initial translation models, one in each  
direction:  $P_{s \rightarrow t}^{(0)}$  and  $P_{t \rightarrow s}^{(0)}$ ;  
3 for  $k=1$  to  $N$  do  
4   Back-translation: Generate source and target  
sentences using the current translation models,  
 $P_{t \rightarrow s}^{(k-1)}$  and  $P_{s \rightarrow t}^{(k-1)}$ , factoring in language  
models,  $P_s$  and  $P_t$ ;  
5   Train new translation models  $P_{s \rightarrow t}^{(k)}$  and  $P_{t \rightarrow s}^{(k)}$   
using the generated sentences and leveraging  $P_s$   
and  $P_t$ ;  
6 end
```



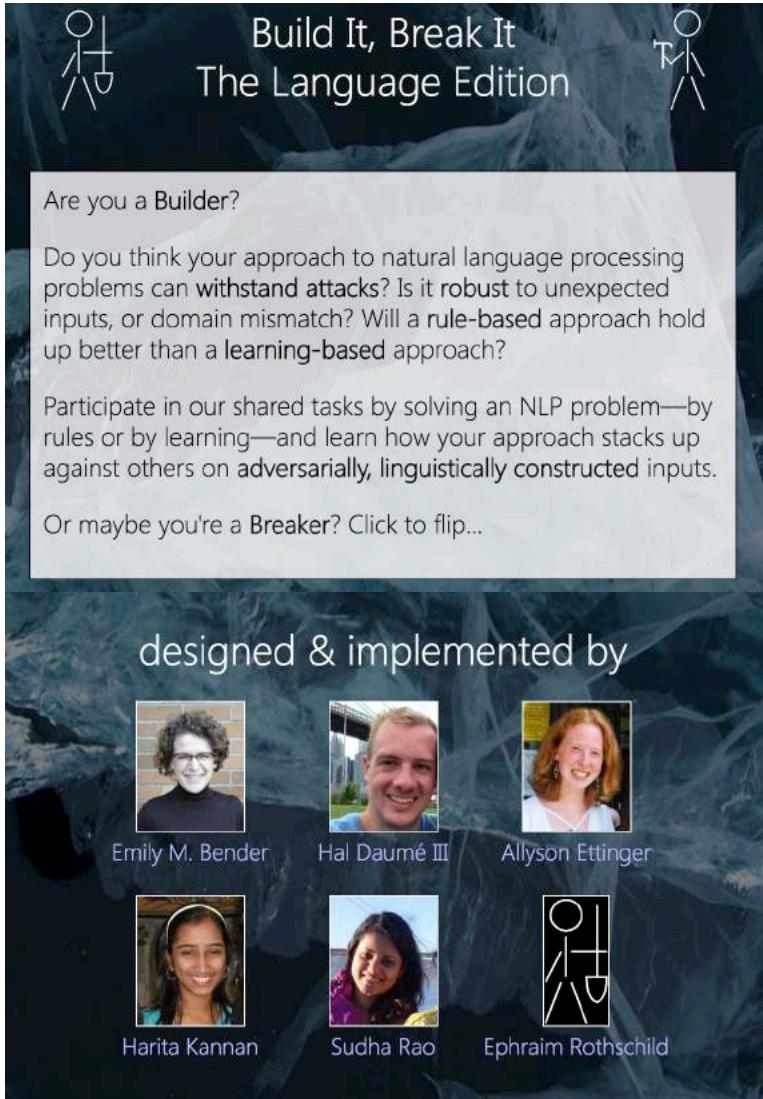
(*Image from Facebook Code Blog)

“Two-dimensional word embeddings in two languages can be aligned via a simple rotation. After the rotation, word translation is performed via nearest neighbor search.”
– [Facebook Code Blog](#)

Unsupervised Machine Translation

- **Unsupervised SMT (IXA NLP Group)**
 - Non-neural: <http://aclweb.org/anthology/D18-1399>
- **Unsupervised NMT (IXA NLP and Cho)**
 - <https://arxiv.org/pdf/1710.11041.pdf>
- **Unsupervised NMT with weights sharing (UCAS)**
 - <http://www.aclweb.org/anthology/P18-1005>
- **Lots of unsupervised MT tech at Facebook at EMNLP 2018**
 - <https://research.fb.com/facebook-research-at-emnlp/>

Robustness and Interpretability



A screenshot of the "Analyzing and interpreting neural networks for NLP" workshop page. The page has a black header with the title in white. Below the title, a dark box contains the text: "Revealing the content of the neural black box: workshop on the analysis and interpretation of neural networks for Natural Language Processing." A red button labeled "View On GitHub" is present. To the right, another dark box states: "This project is maintained by blackboxnlp".

Venue

The workshop will be collocated with EMNLP 2018 in Brussels.

Important dates

- July 19. Submission deadline (11:59pm Pacific Daylight Savings Time, UTC-7h).
- August 3. Notification of acceptance.
- August 30. Camera ready (11:59pm Pacific Daylight Savings Time, UTC-7h).
- November 1. Workshop.

Proceedings

The workshop proceedings are available via ACL Anthology: [proceedings](#)

Workshop program

Time	Program item
09:00-09:10	Opening remarks
09:10-10:00	Invited talk 1: Yoav Goldberg
10:00-11:00	Poster session 1 (10:30-11 tea break)
11:00-12:30	Oral presentation session 1 (6 x 15 minutes)
12:30-14:00	Lunch
14:00-14:50	Invited talk 2: Graham Neubig
14:50-16:00	Poster session 2 (15:30-16 tea break)
16:00-16:50	Invited talk 3: Leila Wehbe
16:50-17:20	Oral presentation session 2 (2 x 15 minutes)



Deep Learning Basics

Recent advancement in Deep Learning in NLP

Perceptron algorithm is a:

*"system that depends on **probabilistic** rather than deterministic principles for its operation, gains its reliability from the **properties of statistical measurements obtain from a large population of elements**"*

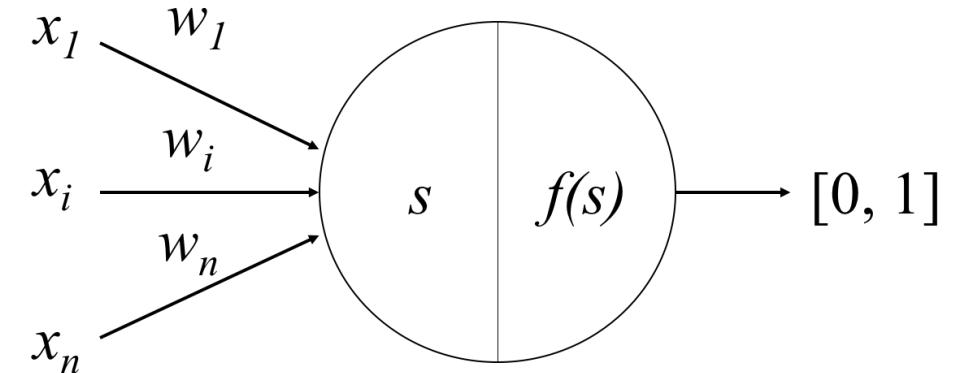
- Frank Rosenblatt (1957)

```
58 # Pick your poison.  
59 from sklearn.linear_model import Perceptron  
60 # Initialize your classifier.  
61 clf = Perceptron(max_iter=10)  
62 # Train the classifier.  
63 clf.fit(X_train, y_train)  
64  
65 print(clf.predict(X_test))
```

Perceptron

Given a **set of inputs x** , perceptron

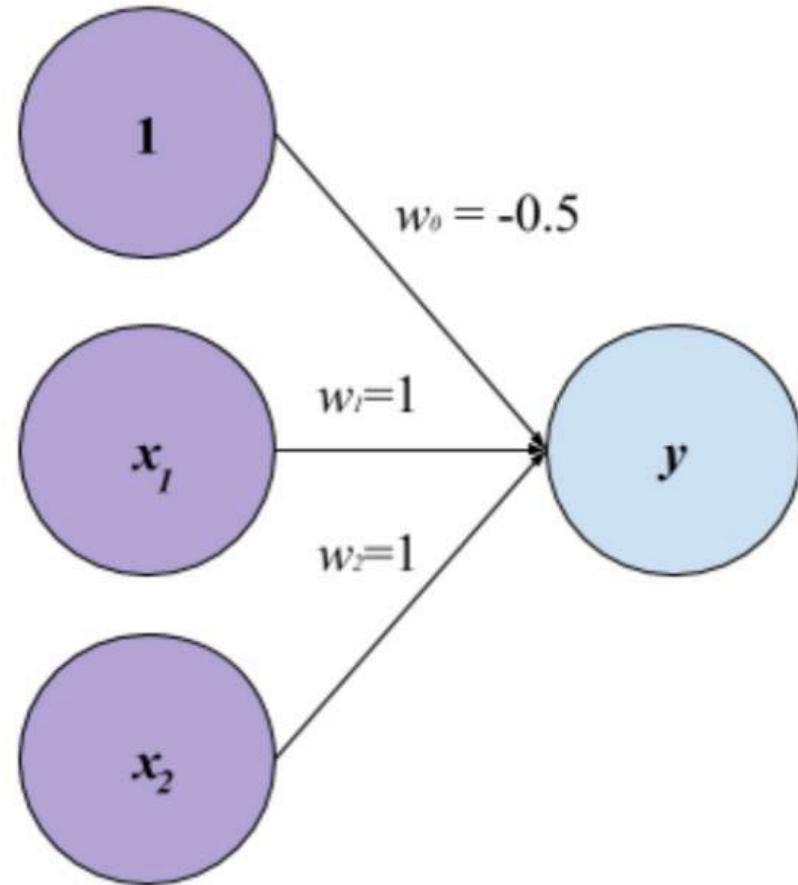
- learns **w vector** to map the inputs to a real-value output between $[0,1]$
- through the **summation of the dot product of the $w \cdot x$**
- with a **transformation function** (aka. activation function)



$$\text{Summation} \\ s = \sum w \cdot x$$

$$\text{Transformation} \\ f(s) = \frac{1}{1+e^{-s}}$$

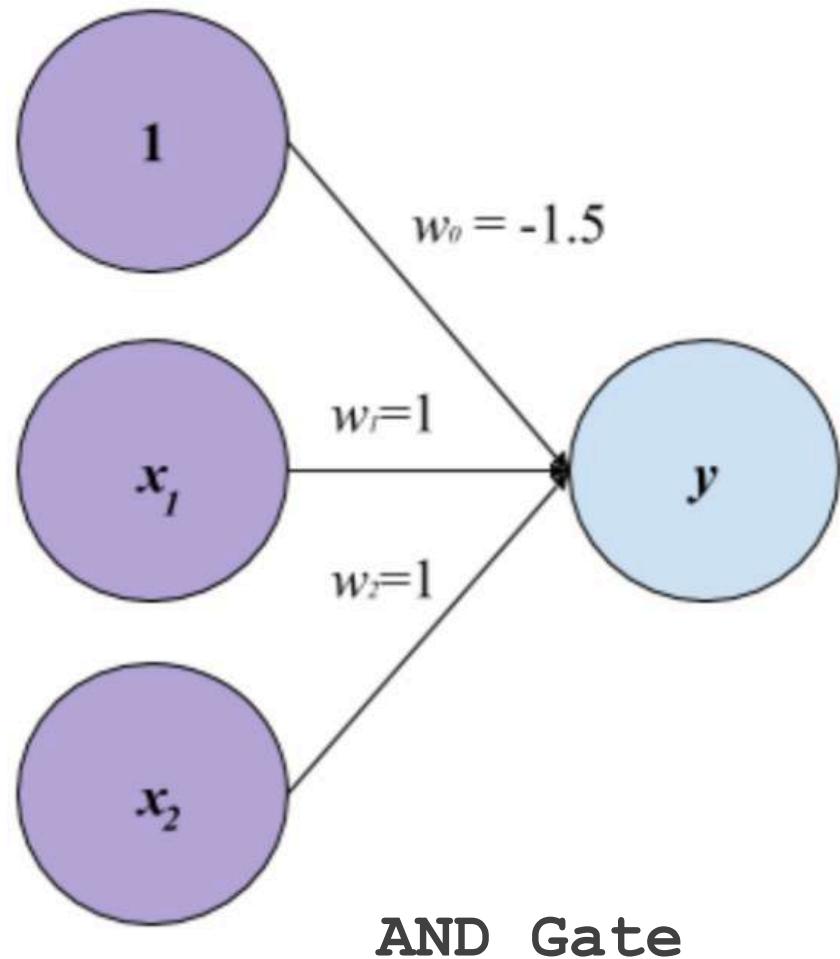
Perceptron



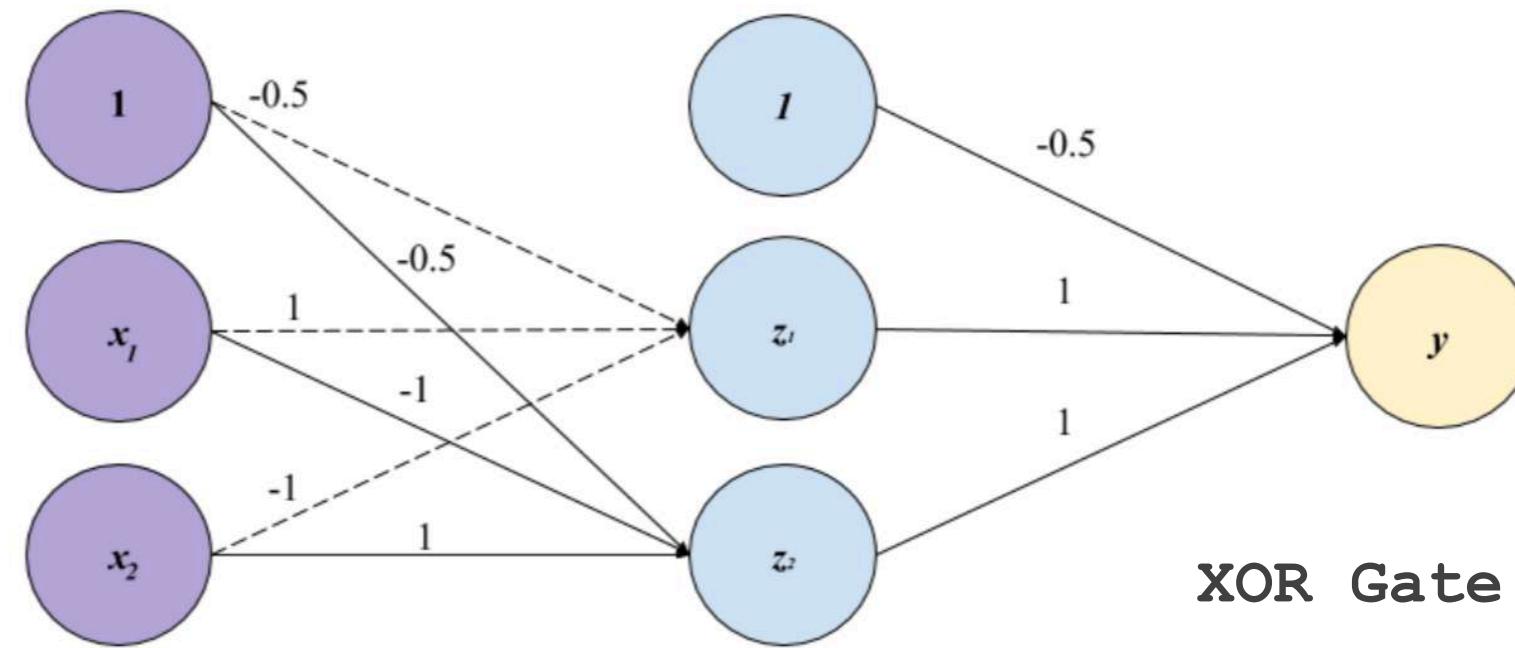
OR Gate

x_1	x_2	$sum = w_0 * 1 + w_1 * x_1 + w_2 * x_2$	$y = 1 \text{ if } sum > 0$ $y = 0 \text{ if } sum \leq 0$
0	0	$-0.5 * 1 + 1 * 0 + 1 * 0 = -0.5$	0
0	1	$-0.5 * 1 + 1 * 0 + 1 * 1 = 0.5$	1
1	0	$-0.5 * 1 + 1 * 1 + 1 * 0 = 0.5$	1
1	1	$-0.5 * 1 + 1 * 1 + 1 * 1 = 1.5$	1

Perceptron



x_1	x_2	$sum = w_0*1 + w_1*x_1 + w_2*x_2$	$y = 1 \text{ if } sum > 0$ $y = 0 \text{ if } sum \leq 0$
0	0	$-1.5*1 + 1*0 + 1*0 = -1.5$	0
0	1	$-1.5*1 + 1*0 + 1*1 = -0.5$	0
1	0	$-1.5*1 + 1*1 + 1*0 = -0.5$	0
1	1	$-1.5*1 + 1*1 + 1*1 = 0.5$	1



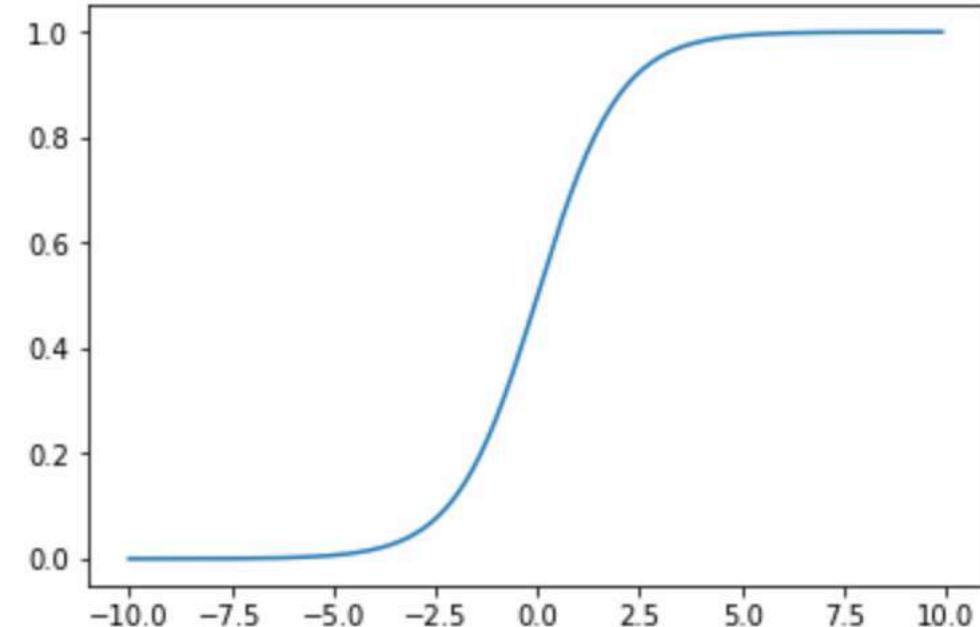
x_1	x_2	$sum_1 = w_0 * I + w_1 * x_1 + w_2 * x_2$	z_1	$sum_2 = w_0 * I + w_1 * x_1 + w_2 * x_2$	z_2
0	0	$-0.5 * 1 + 1 * 0 + -1 * 0 = -0.5$	0	$-0.5 * 1 + -1 * 0 + 1 * 0 = -0.5$	0
0	1	$-0.5 * 1 + 1 * 0 + -1 * 1 = -1.5$	0	$-0.5 * 1 + -1 * 0 + 1 * 1 = 0.5$	1
1	0	$-0.5 * 1 + 1 * 1 + -1 * 0 = 0.5$	1	$-0.5 * 1 + -1 * 1 + 1 * 0 = -1.5$	0
1	1	$-0.5 * 1 + 1 * 1 + -1 * 1 = -0.5$	0	$-0.5 * 1 + -1 * 1 + 1 * 1 = -0.5$	0

z_1	z_2	$sum_1 = w_0 * I + w_1 * x_1 + w_2 * x_2$	y
0	0	$-0.5 * 1 + 1 * 0 + 1 * 0 = -0.5$	0
0	1	$-0.5 * 1 + 1 * 0 + 1 * 1 = 0.5$	1
1	0	$-0.5 * 1 + 1 * 1 + 1 * 0 = 0.5$	1
0	0	$-0.5 * 1 + 1 * 0 + 1 * 0 = -0.5$	0

Figure 2: Emulate an XOR Gate with Feed-forward Network

Activation Function (Sigmoid)

```
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 def sigmoid(x):
7     return 1/(1+np.exp(-x))
8
9 # Generate points from -10 to +10,
10 # in steps of 0.1
11 x = np.arange(-10, 10, 0.1)
12 y = sigmoid(x)
13
14 # Plot the graph.
15 plt.plot(x, y)
16 plt.show()
```



Loss Function (aka. Criterion)

For regression problems, the simplest loss function is to simply take the difference between the predictions and the truth value, i.e. the L1 Loss / Mean Absolute Error (MAE)

Loss Function (aka. Criterion)

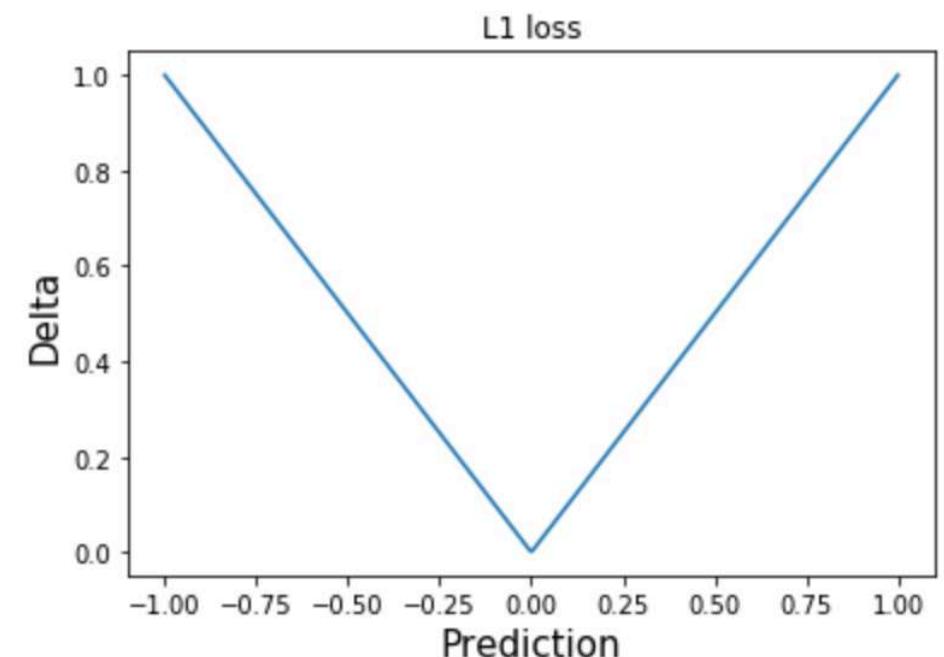
```
38 import numpy as np
39
40 # Create 500 points between -1 and 1.
41 predictions = np.linspace(-1, 1., 500)
42 # Set truth to be the constant 0.
43 truth = np.zeros(500)
44 # Calculate the absolute differences
45 delta = np.abs(truth - predictions)
46
47 # Plotting magic
48 plt.plot(predictions, delta, 'b-', label='L1 loss')
49 plt.title('L1 loss')
50 plt.xlabel('Prediction', fontsize=15)
51 plt.ylabel('Delta', fontsize=15)
```

For regression problems, the simplest loss function is to take the diff between predictions and the truth value, i.e. the L1 Loss / Mean Absolute Error (MAE)

Loss Function (aka. Criterion)

```
38 import numpy as np
39
40 # Create 500 points between -1 and 1.
41 predictions = np.linspace(-1, 1., 500)
42 # Set truth to be the constant 0.
43 truth = np.zeros(500)
44 # Calculate the absolute differences
45 delta = np.abs(truth - predictions)
46
47 # Plotting magic
48 plt.plot(predictions, delta, 'b-', label='L1 loss' )
49 plt.title('L1 loss')
50 plt.xlabel('Prediction', fontsize=15)
51 plt.ylabel('Delta', fontsize=15)
```

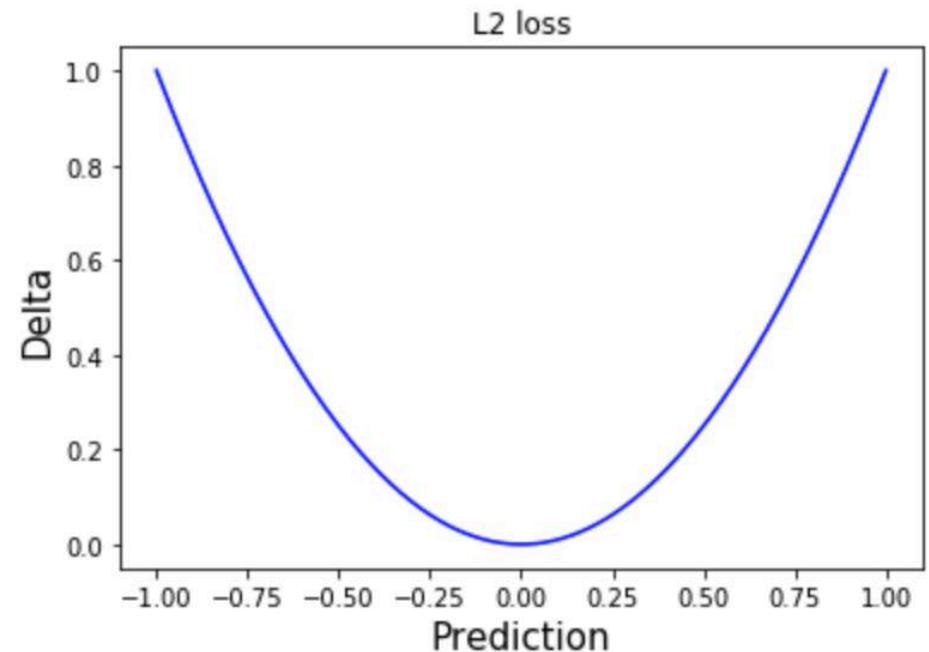
$$L_{mae} = \frac{1}{n} \sum_{i=1}^n |y - \hat{y}|$$



Loss Function (aka. Criterion)

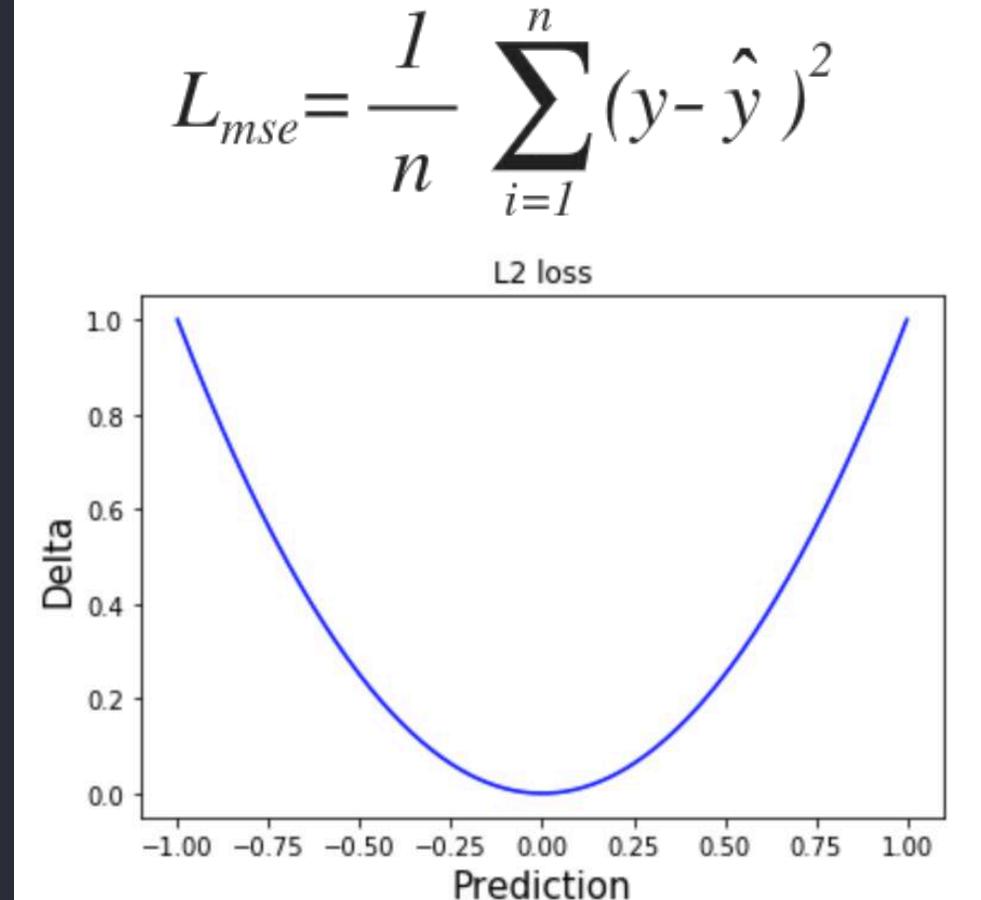
```
38 import numpy as np
39
40 # Create 500 points between -1 and 1.
41 predictions = np.linspace(-1, 1., 500)
42 # Set truth to be the constant 0.
43 truth = np.zeros(500)
44 # Calculate the absolute differences
45 delta = np.abs((truth - predictions)**2) # Line 45
46
47 # Plotting magic
48 plt.plot(predictions, delta, 'b-', label='L2 loss' )
49 plt.title('L2 loss')
50 plt.xlabel('Prediction', fontsize=15)
51 plt.ylabel('Delta', fontsize=15)
```

$$L_{mse} = \frac{1}{n} \sum_{i=1}^n (y - \hat{y})^2$$

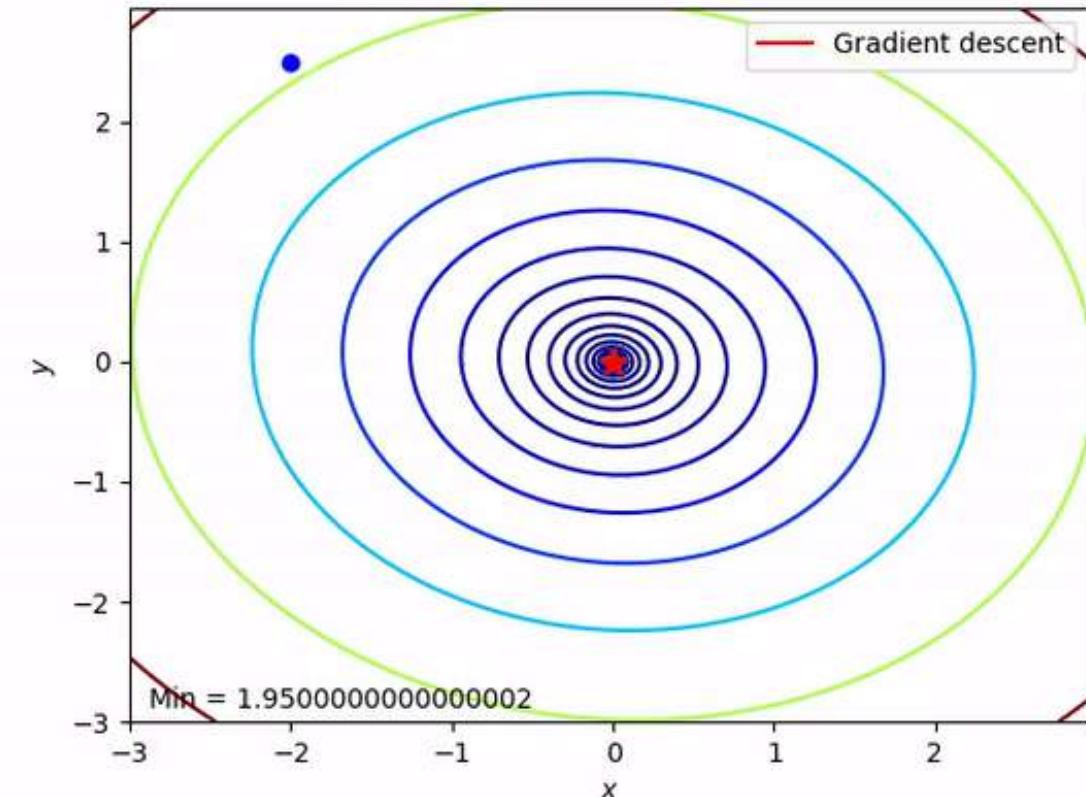
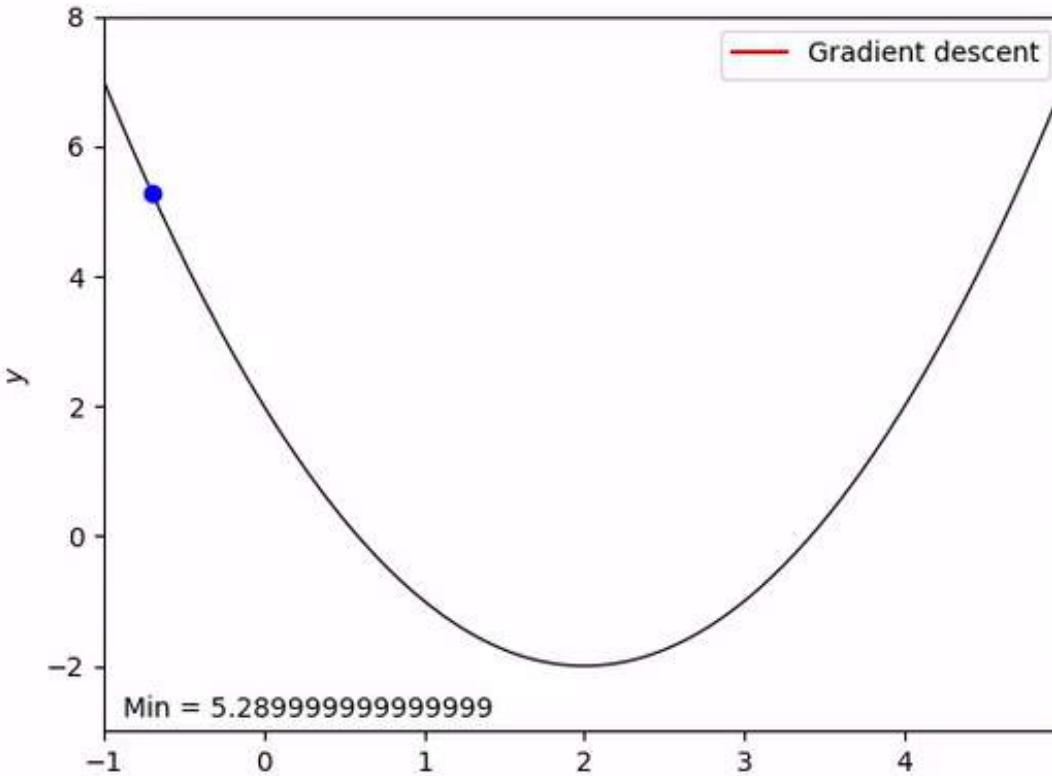


Loss Function (aka. Criterion)

```
38 import numpy as np
39
40 # Create 500 points between -1 and 1.
41 predictions = np.linspace(-1, 1., 500)
42 # Set truth to be the constant 0.
43 truth = np.zeros(500)
44 # Calculate the absolute differences
45 delta = np.abs((truth - predictions)**2)
46
47 # Plotting magic
48 plt.plot(predictions, delta, 'b-', label='L2 loss' )
49 plt.title('L2 loss')
50 plt.xlabel('Prediction', fontsize=15)
51 plt.ylabel('Delta', fontsize=15)
```



Optimization (Gradient Descent)



(Images from https://jed-ai.github.io/py1_gd_animation/)
It has some cool code to generate the GD pictures

Typically, process performs the following 4 steps iteratively.

Initialization

1. Initialize weights vector

Forward Propagation

- 2a. Multiply the weights vector with the inputs, sum the products.
2b. Put the sum through the activation function, e.g. sigmoid

Back Propagation

- 3a. Compute the errors, i.e. difference between expected output and predictions
- 3b. Multiply the error with the derivatives to get the delta
- 3c. Multiply the delta vector with the inputs, sum the product

Optimizer takes a step

4. Multiply the learning rate with the output of step 3c

Repeat 1-4 until desired



Summary

- **Classic NLP == Lots of feature engineering**
 - Frequency, TF-IDF, PPMI
- **Deep ‘Magic’ NLP**
 - Transfer Learning moment for NLP
 - NMT achieving human parity & Unsupervised MT
 - Robustness and Interpretability
- **Deep Learning Basics**
 - Perceptron and Multi-layered Perceptron
 - Activation function, Loss function, Gradient Descent
 - Training routine

Environment Setup

Open Anaconda Navigator.

Go to the PyTorch installation page, copy the command as per configuration:

<https://pytorch.org/get-started/locally/>

Fire up the terminal in Anaconda Navigator.

Start a Jupyter Notebook.

Download <https://goo.gl/fyZv3m>

Import the .ipynb to the Jupyter Notebook

Fin



Neural Nets from Scratch

(with some numpy magic)

```
In [1]: import sys  
sys.executable
```

```
Out[1]: '/usr/local/bin/python3'
```

```
In [2]: !which pip3  
!which python3  
!ls -lah /usr/local/bin/python3
```

```
/Library/Frameworks/Python.framework/Versions/3.6/bin/pip3  
/Library/Frameworks/Python.framework/Versions/3.6/bin/python3  
lrwxr-xr-x 1 liling.tan admin 69B Dec 14 2017 /usr/local/bin/python3 -> ../../../../../../Library/Frameworks/Python.framework/Versions/3.6/bin/python3
```

```
In [3]: !pip3 install torch==1.0.0
```

```
Requirement already satisfied: torch==1.0.0 in /Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages (1.0.0)
```

```
In [4]: import torch  
torch.cuda.is_available()
```

```
Out[4]: False
```

These are just some IPython functions for showing image and making things pretty



```
In [1]: # IPython candies...
from IPython.display import Image
from IPython.core.display import HTML

from IPython.display import clear_output
```

```
In [2]: %%html
<style> table {float:left} </style>
```

Perceptron

Perceptron algorithm is a:

"*system that depends on **probabilistic** rather than deterministic principles for its operation, gains its reliability from the **properties of statistical measurements obtain from a large population of elements***" - Frank Rosenblatt (1957)

Then the news:

"[Perceptron is an] **embryo of an electronic computer** that [the Navy] expects will be **able to walk, talk, see, write, reproduce itself and be conscious of its existence.**" - The New York Times (1958)

News quote cite from Olazaran (1996)

Perceptron in Bullets

- Perceptron learns to classify any linearly separable set of inputs.
- Some nice graphics for perceptron with Go <https://appliedgo.net/perceptron/>

If you've got some spare time:

- There's a whole book just on perceptron: <https://mitpress.mit.edu/books/perceptrons>
- For watercooler gossips on perceptron in the early days, read [Olazaran \(1996\)](#)

Perceptron in Math

Given a set of inputs x , the perceptron

- learns w vector to map the inputs to a real-value output between $[0, 1]$
- through the summation of the dot product of the $w \cdot x$ with a transformation function

Perceptron as a Workflow Diagram

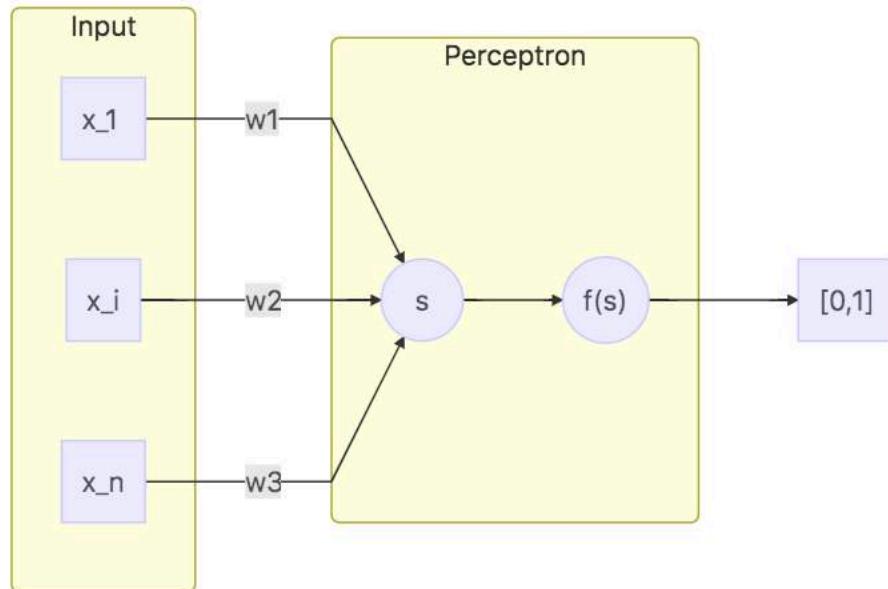
If you're familiar with [mermaid flowchart](#)

```
.. mermaid:::
```

```
graph LR
    x_1 -->|w1| n1((s))
    x_i -->|w2| n1((s))
    x_n -->|w3| n1((s))
    subgraph Perceptron
        n1 --> n2(("f(s)"))
    end
    n2 --> y["[0,1]"]
```

```
In [6]: ##Image(url="perceptron-mermaid.svg", width=500)
Image(url="https://svgshare.com/i/AbJ.svg", width=500)
```

Out[6]:



Optimization Process

To learn the weights, w , we use an **optimizer** to find the best-fit (optimal) values for w such that the inputs correctly map to the outputs.

Typically, process performs the following 4 steps iteratively.

Initialization

- **Step 1:** Initialize weights vector

Forward Propagation

- **Step 2a:** Multiply the weights vector with the inputs, sum the products, i.e. s
- **Step 2b:** Put the sum through the sigmoid, i.e. $f()$

Back Propagation

- **Step 3a:** Compute the errors, i.e. difference between expected output and predictions
- **Step 3b:** Multiply the error with the **derivatives** to get the delta
- **Step 3c:** Multiply the delta vector with the inputs, sum the product

Optimizer takes a step

- **Step 4:** Multiply the learning rate with the output of Step 3c.

```
In [62]: import math
import numpy as np
np.random.seed(0)

def sigmoid(x): # Returns values that sums to one.
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(sx):
    # See https://math.stackexchange.com/a/1225116
    # Hint: let sx = sigmoid(x)
    return ???
```

Fill in the blanks =)

```
In [9]: sigmoid(np.array([2.5, 0.32, -1.42])) # [out]: array([0.92414182, 0.57932425, 0.19466158])
```

```
Out[9]: array([0.92414182, 0.57932425, 0.19466158])
```

```
In [10]: sigmoid_derivative(np.array([2.5, 0.32, -1.42])) # [out]: array([0.07010372, 0.24370766, 0.15676845])
```

```
Out[10]: array([0.07010372, 0.24370766, 0.15676845])
```

```
In [62]: import math
import numpy as np
np.random.seed(0)
```

```
def sigmoid(x): # Returns values that sums to one.
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(sx):
    # See https://math.stackexchange.com/a/1225116
    # Hint: let sx = sigmoid(x)
    return sx * (1 -sx)
```

```
In [9]: sigmoid(np.array([2.5, 0.32, -1.42]))          # [out]: array([0.92414182, 0.57932425, 0.19466158])
```

```
Out[9]: array([0.92414182, 0.57932425, 0.19466158])
```

```
In [10]: sigmoid_derivative(np.array([2.5, 0.32, -1.42])) # [out]: array([0.07010372, 0.24370766, 0.15676845])
```

```
Out[10]: array([0.07010372, 0.24370766, 0.15676845])
```

```
def cost(predicted, truth):  
    return np.abs(truth - predicted)
```

```
gold = np.array([0.5, 1.2, 9.8])  
pred = np.array([0.6, 1.0, 10.0])  
cost(pred, gold)
```

```
array([0.1, 0.2, 0.2])
```

```
gold = np.array([0.5, 1.2, 9.8])  
pred = np.array([9.3, 4.0, 99.0])  
cost(pred, gold)
```

```
array([ 8.8,  2.8, 89.2])
```

Representing OR Boolean

Lets consider the problem of the OR boolean and apply the perceptron with simple gradient descent.

x2	x3	y
0	0	0
0	1	1
1	0	1
1	1	1

```
X = or_input = np.array([[0,0], [0,1], [1,0], [1,1]])
Y = or_output = np.array([[0,1,1,1]]).T
```

`or_input`

```
array([[0, 0],
       [0, 1],
       [1, 0],
       [1, 1]])
```

`or_output`

```
array([[0],
       [1],
       [1],
       [1]])
```

```
# Define the shape of the weight vector.  
num_data, input_dim = ???  
# Define the shape of the output vector.  
output_dim = len(or_output.T)
```

Fill in the blanks =)

```
print('Inputs\n=====')  
print('no. of rows =', num_data)  
print('no. of cols =', input_dim)  
print('\n')  
print('Outputs\n=====')  
print('no. of cols =', output_dim)
```

Inputs
=====

no. of rows = 4
no. of cols = 2

Outputs
=====

no. of cols = 1

```
# Initialize weights between the input layers and the perceptron  
W = np.random.random((input_dim, output_dim))  
W  
  
array([[0.5488135 ,  
       [0.71518937]])
```

```
# Define the shape of the weight vector.  
num_data, input_dim = or_input.shape  
# Define the shape of the output vector.  
output_dim = len(or_output.T)
```

```
print('Inputs\n=====')  
print('no. of rows =', num_data)  
print('no. of cols =', input_dim)  
print('\n')  
print('Outputs\n=====')  
print('no. of cols =', output_dim)
```

Inputs
=====

```
no. of rows = 4  
no. of cols = 2
```

Outputs
=====

```
no. of cols = 1
```

```
# Initialize weights between the input layers and the perceptron  
W = np.random.random((input_dim, output_dim))  
W  
  
array([[0.5488135 ,  
       [0.71518937]])
```

Step 2a: Multiply the weights vector with the inputs, sum the products

To get the output of step 2a,

- Iterate through each row of the data, x
- For each column in each row, find the product of the value and the respective weights
- For each row, compute the sum of the products

```
# If we write it imperatively:  
summation = []  
for row in X:  
    sum_wx = 0  
    for feature, weight in zip(row, w):  
        sum_wx += ???  
    summation.append(???)  
print(np.array(summation))
```

Fill in the blanks =)

```
[[0.      ]  
 [0.38344152]  
 [0.96366276]  
 [1.34710428]]
```

Step 2a: Multiply the weights vector with the inputs, sum the products

To get the output of step 2a,

- Iterate through each row of the data, x
- For each column in each row, find the product of the value and the respective weights
- For each row, compute the sum of the products

```
# If we write it imperatively:  
summation = []  
for row in X:  
    sum_wx = 0  
    for feature, weight in zip(row, w):  
        sum_wx += feature * weight  
    summation.append(sum_wx)  
print(np.array(summation))
```

```
[[0.  
     ]  
 [0.38344152]  
 [0.96366276]  
 [1.34710428]]
```

Step 2a: Multiply the weights vector with the inputs, sum the products

To get the output of step 2a,

- Iterate through each row of the data, x
- For each column in each row, find the product of the value and the respective weights
- For each row, compute the sum of the products

```
# If we write it imperatively:
summation = []
for row in X:
    sum_wx = 0
    for feature, weight in zip(row, w):
        sum_wx += feature * weight
    summation.append(sum_wx)
print(np.array(summation))
```

```
[[0.      ]
 [0.38344152]
 [0.96366276]
 [1.34710428]]
```

```
# If we vectorize the process and use numpy.
np.dot(X, w)
```

```
array([[0.      ],
 [0.38344152],
 [0.96366276],
 [1.34710428]])
```

Train the Single-Layer Model

```
num_epochs = 10000 # No. of times to iterate.
learning_rate = 0.03 # How large a step to take per iteration.

# Lets standardize and call our inputs X and outputs Y
X = or_input
Y = or_output

for _ in range(num_epochs):
    layer0 = X

    # Step 2a: Multiply the weights vector with the inputs, sum the products, i.e. s
    # Step 2b: Put the sum through the sigmoid, i.e. f()
    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(X, W))

    # Back propagation.
    # Step 3a: Compute the errors, i.e. difference between expected output and predictions
    # How much did we miss?
    layer1_error = cost(layer1, Y)

    # Step 3b: Multiply the error with the derivatives to get the delta
    # multiply how much we missed by the slope of the sigmoid at the values in layer1
    layer1_delta = layer1_error * sigmoid_derivative(layer1)

    # Step 3c: Multiply the delta vector with the inputs, sum the product (use np.dot)
    # Step 4: Multiply the learning rate with the output of Step 3c.
    W += learning_rate * np.dot(layer0.T, layer1_delta)
```

```
layer1
```

```
array([[0.5  
       ],  
      [0.95603077],  
      [0.95677329],  
      [0.99792643]])
```

```
# Expected output.
```

```
Y
```

```
array([[0],  
      [1],  
      [1],  
      [1]])
```

```
# On the training data
```

```
[[int(prediction > 0.5)] for prediction in layer1]
```

```
[[0], [1], [1], [1]]
```

Lets try the XOR Boolean

Lets consider the problem of the OR boolean and apply the perceptron with simple gradient descent.

x2	x3	y
0	0	0
0	1	1
1	0	1
1	1	0

```
X = xor_input = np.array([[0,0], [0,1], [1,0], [1,1]])  
Y = xor_output = np.array([[0,1,1,0]]).T
```

```
xor_input
```

```
array([[0, 0],  
       [0, 1],  
       [1, 0],  
       [1, 1]])
```

```
xor_output
```

```
array([[0],  
       [1],  
       [1],  
       [0]])
```

```

num_epochs = 10000 # No. of times to iterate.
learning_rate = 0.003 # How large a step to take per iteration.

# Lets drop the last row of data and use that as unseen test.
X = xor_input
Y = xor_output

# Define the shape of the weight vector.
num_data, input_dim = ???
# Define the shape of the output vector.
output_dim = len(Y.T)
# Initialize weights between the input layers and the perceptron
W = np.random.random((input_dim, output_dim))

for _ in range(num_epochs):
    layer0 = X
    # Forward propagation.
    # Inside the perceptron, Step 2.
    layer1 = ???

    # How much did we miss?
    layer1_error = ???

    # Back propagation.
    # multiply how much we missed by the slope of the sigmoid at the values in layer1
    layer1_delta = ???

    # update weights
    W += ???

```

Fill in the blanks =)

```
num_epochs = 10000 # No. of times to iterate.
learning_rate = 0.003 # How large a step to take per iteration.

# Lets drop the last row of data and use that as unseen test.
X = xor_input
Y = xor_output

# Define the shape of the weight vector.
num_data, input_dim = X.shape
# Define the shape of the output vector.
output_dim = len(Y.T)
# Initialize weights between the input layers and the perceptron
W = np.random.random((input_dim, output_dim))

for _ in range(num_epochs):
    layer0 = X
    # Forward propagation.
    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(X, W))

    # How much did we miss?
    layer1_error = cost(layer1, Y)

    # Back propagation.
    # multiply how much we missed by the slope of the sigmoid at the values in layer1
    layer1_delta = sigmoid_derivative(layer1) * layer1_error

    # update weights
    W += learning_rate * np.dot(layer0.T, layer1_delta)
```

```
# Expected output.
```

```
Y
```

```
array([[0],  
       [1],  
       [1],  
       [0]])
```

```
# On the training data
```

```
[int(prediction > 0.5) for prediction in layer1] # All correct.
```

```
[0, 1, 1, 1]
```

Wrong!!!

You can't represent XOR with simple perceptron !!!

No matter how you change the hyperparameters or data, the XOR function can't be represented by a single perceptron layer.

There's no way you can get all four data points to get the correct outputs for the XOR boolean operation.

Solving XOR (Add more layers)

```
from itertools import chain
import numpy as np
np.random.seed(0)

def sigmoid(x): # Returns values that sums to one.
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(sx):
    # See https://math.stackexchange.com/a/1225116
    return sx * (1 - sx)

# Cost functions.
def cost(predicted, truth):
    return truth - predicted

xor_input = np.array([[0,0], [0,1], [1,0], [1,1]])
xor_output = np.array([[0,1,1,0]]).T

# Lets drop the last row of data and use that as unseen test.
X = xor_input[:-1]
Y = xor_output[:-1]

# Define the shape of the weight vector.
num_data, input_dim = X.shape
# Lets set the dimensions for the intermediate layer.
hidden_dim = 5
```

```

hidden_dim = 5
# Initialize weights between the input layers and the hidden layer.
W1 = np.random.random((input_dim, hidden_dim))

# Define the shape of the output vector.
output_dim = len(Y.T)
# Initialize weights between the hidden layers and the output layer.
W2 = ???

# Initialize weigh
num_epochs = 10000
learning_rate = 0.03

for epoch_n in range(num_epochs):
    layer0 = X
    # Forward propagation.

    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(layer0, W1))
    layer2 = ???

    # Back propagation (Y -> layer2)

    # How much did we miss in the predictions?
    layer2_error = cost(layer2, Y)
    # In what direction is the target value?
    # Were we really close? If so, don't change too much.
    layer2_delta = layer2_error * sigmoid_derivative(layer2)

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer2 error (according to the weights)?
    layer1_error = np.dot(layer2_delta, W2.T)
    layer1_delta = ???

    # update weights
    W2 += learning_rate * np.dot(layer1.T, layer2_delta)
    W1 += ???

```

Fill in the blanks =)

```
hidden_dim = 5
# Initialize weights between the input layers and the hidden layer.
W1 = np.random.random((input_dim, hidden_dim))

# Define the shape of the output vector.
output_dim = len(Y.T)
# Initialize weights between the hidden layers and the output layer.
W2 = np.random.random((hidden_dim, output_dim))

# Initialize weigh
num_epochs = 10000
learning_rate = 0.03

for epoch_n in range(num_epochs):
    layer0 = X
    # Forward propagation.

    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(layer0, W1))
    layer2 = sigmoid(np.dot(layer1, W2))

    # Back propagation (Y -> layer2)

    # How much did we miss in the predictions?
    layer2_error = cost(layer2, Y)
    # In what direction is the target value?
    # Were we really close? If so, don't change too much.
    layer2_delta = layer2_error * sigmoid_derivative(layer2)

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer2 error (according to the weights)?
    layer1_error = np.dot(layer2_delta, W2.T)
    layer1_delta = layer1_error * sigmoid_derivative(layer1)

    # update weights
    W2 += learning_rate * np.dot(layer1.T, layer2_delta)
    W1 += learning_rate * np.dot(layer0.T, layer1_delta)
```

```
# Training input.
```

```
X
```

```
array([[1, 1],  
       [0, 1],  
       [0, 0],  
       [1, 0]])
```

```
# Expected output.
```

```
Y
```

```
array([[0],  
       [1],  
       [0],  
       [1]])
```

```
layer2 # Our output layer
```

```
array([[0.46427804],  
       [0.6213127 ],  
       [0.31284349],  
       [0.62323891]])
```

```
# On the training data
```

```
[int(prediction > 0.5) for prediction in layer2]
```

```
[0, 1, 0, 1]
```

Now try adding another layer

Use the same process:

1. Initialize
2. Forward Propagate
3. Back Propagate
4. Update (aka step)

```
from itertools import chain
import numpy as np
np.random.seed(0)

def sigmoid(x): # Returns values that sums to one.
    return 1 / (1 + np.exp(-x))

def sigmoid_derivative(sx):
    # See https://math.stackexchange.com/a/1225116
    return sx * (1 - sx)

# Cost functions.
def cost(predicted, truth):
    return truth - predicted

xor_input = np.array([[0,0], [0,1], [1,0], [1,1]])
xor_output = np.array([[0,1,1,0]]).T

# Lets drop the last row of data and use that as unseen test.
X = xor_input_shuff
Y = xor_output_shuff
```

Fill in the blanks =)

```
layer0to1_hidden_dim = 5
layer1to2_hidden_dim = 5

# Initialize weights between the input layers 0 -> layer 1
W1 = np.random.random((input_dim, ???))

# Initialize weights between the layer 1 -> layer 2
W2 = np.random.random(???, ???)

# Define the shape of the output vector.
output_dim = len(Y.T)
# Initialize weights between the layer 2 -> layer 3
W3 = np.random.random(???, output_dim)

# Initialize weigh
num_epochs = 10000
learning_rate = 1.0

for epoch_n in range(num_epochs):
    layer0 = X
    # Forward propagation.

    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(layer0, W1))
    layer2 = ???
    layer3 = ???

    # Back propagation (Y -> layer2)
    # How much did we miss in the predictions?
    layer3_error = ???
    # In what direction is the target value?
    # Were we really close? If so, don't change too much.
    layer3_delta = ???

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer3 error (according to the weights)?
    layer2_error = ???
    layer2_delta = ???

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer2 error (according to the weights)?
    layer1_error = ???
    layer1_delta = layer1_error * sigmoid_derivative(layer1)

    # update weights
    W3 += learning_rate * np.dot(layer2.T, layer3_delta)
    W2 += ???
    W1 += ???
```

```

layer0to1_hidden_dim = 5
layer1to2_hidden_dim = 5

# Initialize weights between the input layers 0 -> layer 1
W1 = np.random.random((input_dim, layer0to1_hidden_dim))

# Initialize weights between the layer 1 -> layer 2
W2 = np.random.random((layer0to1_hidden_dim, layer1to2_hidden_dim))

# Define the shape of the output vector.
output_dim = len(Y.T)
# Initialize weights between the layer 2 -> layer 3
W3 = np.random.random((layer1to2_hidden_dim, output_dim))

# Initialize weigh
num_epochs = 10000
learning_rate = 1.0

for epoch_n in range(num_epochs):
    layer0 = X
    # Forward propagation.

    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(layer0, W1))
    layer2 = sigmoid(np.dot(layer1, W2))
    layer3 = sigmoid(np.dot(layer2, W3))

    # Back propagation (Y -> layer2)
    # How much did we miss in the predictions?
    layer3_error = cost(layer3, Y)
    # In what direction is the target value?
    # Were we really close? If so, don't change too much.
    layer3_delta = layer3_error * sigmoid_derivative(layer3)

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer3 error (according to the weights)?
    layer2_error = np.dot(layer3_delta, W3.T)
    layer2_delta = layer3_error * sigmoid_derivative(layer2)

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer2 error (according to the weights)?
    layer1_error = np.dot(layer2_delta, W2.T)
    layer1_delta = layer1_error * sigmoid_derivative(layer1)

    # update weights
    W3 += learning_rate * np.dot(layer2.T, layer3_delta)
    W2 += learning_rate * np.dot(layer1.T, layer2_delta)
    W1 += learning_rate * np.dot(layer0.T, layer1_delta)

```

Now, lets do it with PyTorch

First lets try a single perceptron and see that we can't train a model that can represent XOR.

```
from tqdm import tqdm

import torch
from torch import nn
from torch.autograd import Variable
from torch import FloatTensor
from torch import optim
use_cuda = torch.cuda.is_available()
```

```
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

sns.set_style("darkgrid")
sns.set(rc={'figure.figsize':(15, 10)})
```

```
X # Original XOR X input in numpy array data structure.
```

```
array([[0, 1],  
       [1, 1],  
       [1, 0],  
       [0, 0]])
```

```
Y # Original XOR Y output in numpy array data structure.
```

```
array([[1],  
       [0],  
       [1],  
       [0]])
```

```
device = 'gpu' if torch.cuda.is_available() else 'cpu'  
# Converting the X to PyTorch-able data structure.  
X_pt = torch.tensor(X).float()  
X_pt = X_pt.to(device)  
# Converting the Y to PyTorch-able data structure.  
Y_pt = torch.tensor(Y, requires_grad=False).float()  
Y_pt = Y_pt.to(device)  
print(X_pt)  
print(Y_pt)
```

```
tensor([[0., 1.],  
       [1., 1.],  
       [1., 0.],  
       [0., 0.]])
```

```
tensor([[1.],  
       [0.],  
       [1.],  
       [0.]])
```

```
# Use tensor.shape to get the shape of the matrix/tensor.  
num_data, input_dim = X_pt.shape  
print('Inputs Dim:', input_dim)  
  
num_data, output_dim = Y_pt.shape  
print('Output Dim:', output_dim)
```

```
Inputs Dim: 2  
Output Dim: 1
```

```
# Use Sequential to define a simple feed-forward network.  
model = nn.Sequential(  
    nn.Linear(input_dim, output_dim), # Use nn.Linear to get our simple perceptron  
    nn.Sigmoid() # Use nn.Sigmoid to get our sigmoid non-linearity  
)  
model
```

```
Sequential(  
    (0): Linear(in_features=2, out_features=1, bias=True)  
    (1): Sigmoid()  
)
```

```
# Remember we define as: cost = truth - predicted  
# If we take the absolute of cost, i.e.: cost = |truth - predicted|  
# we get the L1 loss function.  
criterion = nn.L1Loss()  
learning_rate = 0.03
```

```
# The simple weights/parameters update processes we did before  
# is call the gradient descent. SGD is the stochastic variant of  
# gradient descent.  
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
```

Training a PyTorch model

To train a model using PyTorch, we simply iterate through the no. of epochs and imperatively state the computations we want to perform.

Remember the steps?

1. Initialize
2. Forward Propagation
3. Backward Propagation
4. Update Optimizer

```
num_epochs = 7000

# Step 1: Initialization.
# Note: When using PyTorch a lot of the manual weights
#       initialization is done automatically when we define
#       the model (aka architecture)
model = nn.Sequential(
    nn.Linear(input_dim, output_dim),
    nn.Sigmoid())
criterion = nn.L1Loss()
learning_rate = 1.0
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
num_epochs = 10000

losses = []
```

```

# Step 1: Initialization.
# Note: When using PyTorch a lot of the manual weights
#       initialization is done automatically when we define
#       the model (aka architecture)
model = nn.Sequential(
    nn.Linear(input_dim, output_dim),
    nn.Sigmoid())
criterion = nn.MSELoss()
learning_rate = 1.0
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
num_epochs = 10000

losses = []

for i in tqdm(range(num_epochs)):
    # Reset the gradient after every epoch.
    optimizer.zero_grad()
    # Step 2: Foward Propagation
    predictions = model(X_pt)

    # Step 3: Back Propagation
    # Calculate the cost between the predictions and the truth.
    loss_this_epoch = criterion(predictions, Y_pt)
    # Note: The neat thing about PyTorch is it does the
    #       auto-gradient computation, no more manually defining
    #       derivative of functions and manually propagating
    #       the errors layer by layer.
    loss_this_epoch.backward()

    # Step 4: Optimizer take a step.
    # Note: Previously, we have to manually update the
    #       weights of each layer individually according to the
    #       learning rate and the layer delta.
    #       PyTorch does that automatically =
    optimizer.step()

    # Log the loss value as we proceed through the epochs.
    losses.append(loss_this_epoch.data.item())

# Visualize the losses
plt.plot(losses)
plt.show()

```

```
for _x, _y in zip(X_pt, Y_pt):
    prediction = model(_x)
    print('Input:\t', list(map(int, _x)))
    print('Pred:\t', int(prediction))
    print('Output:\t', int(_y))
    print('#####')
```

Input: [0, 1]

Pred: 0

Ouput: 1

#####

Input: [1, 1]

Pred: 1

Ouput: 0

#####

Input: [0, 0]

Pred: 0

Ouput: 0

#####

Input: [1, 0]

Pred: 0

Ouput: 1

#####

Now, try again with 2 layers using PyTorch

```

%%time

hidden_dim = 5
num_data, input_dim = X_pt.shape
num_data, output_dim = Y_pt.shape

model = nn.Sequential(nn.Linear(???, ??),
                      ??,
                      nn.Linear(???, ??),
                      nn.Sigmoid())

criterion = nn.MSELoss()
learning_rate = 0.3
optimizer = ???
num_epochs = 5000

losses = []

for _ in tqdm(range(num_epochs)):
    optimizer.zero_grad()
    predictions = ???
    loss_this_epoch = ???
    loss_this_epoch.???
    optimizer.???
    losses.append(loss_this_epoch.data.item())
    ##print([float(_pred) for _pred in predictions], list(map(int, Y_pt)), loss_this_epoch.data[0])

# Visualize the losses
plt.plot(losses)
plt.show()

```

Fill in the blanks =)

Now, try again with 2 layers using PyTorch

```
%time

hidden_dim = 5
num_data, input_dim = X_pt.shape
num_data, output_dim = Y_pt.shape

model = nn.Sequential(nn.Linear(input_dim, hidden_dim),
                      nn.Sigmoid(),
                      nn.Linear(hidden_dim, output_dim),
                      nn.Sigmoid())

criterion = nn.MSELoss()
learning_rate = 0.3
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
num_epochs = 5000

losses = []

for _ in tqdm(range(num_epochs)):
    optimizer.zero_grad()
    predictions = model(X_pt)
    loss_this_epoch = criterion(predictions, Y_pt)
    loss_this_epoch.backward()
    optimizer.step()
    losses.append(loss_this_epoch.data.item())
    ##print([float(_pred) for _pred in predictions], list(map(int, Y_pt)), loss_this_epoch.data[0])

# Visualize the losses
plt.plot(losses)
plt.show()
```

```
: for _x, _y in zip(X_pt, Y_pt):
    prediction = model(_x)
    print('Input:\t', list(map(int, _x)))
    print('Pred:\t', int(prediction > 0.5))
    print('Ouput:\t', int(_y))
    print('#####')
```

Input: [1, 1]

Pred: 0

Ouput: 0

#####

Input: [0, 1]

Pred: 1

Ouput: 1

#####

Input: [0, 0]

Pred: 0

Ouput: 0

#####

Input: [1, 0]

Pred: 1

Ouput: 1

#####



MNIST

The “Hello World” of Neural Nets

MNIST: The "Hello World" of Neural Nets

Like any deep learning class, we **must** do the MNIST.

The MNIST dataset is

- is made up of handwritten digits
- 60,000 examples training set
- 10,000 examples test set

```
# We're going to install tensorflow here because their dataset access is simpler =)
!pip3 install tensorflow
```

```
Requirement already satisfied: tensorflow in /Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages (1.8.0)
Requirement already satisfied: termcolor>=1.1.0 in /Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages (from tensorflow) (1.1.0)
```

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

```
# Visualization Candies
import matplotlib.pyplot as plt

def show_image(mnist_x_vector, mnist_y_vector):
    pixels = mnist_x_vector.reshape((28, 28))
    label = np.where(mnist_y_vector == 1)[0][0]
    plt.title('Label is {}'.format(label))
    plt.imshow(pixels, cmap='gray')
    plt.show()
```

```
# Fifth image and label.
show_image(mnist.train.images[5], mnist.train.labels[5])
```

```

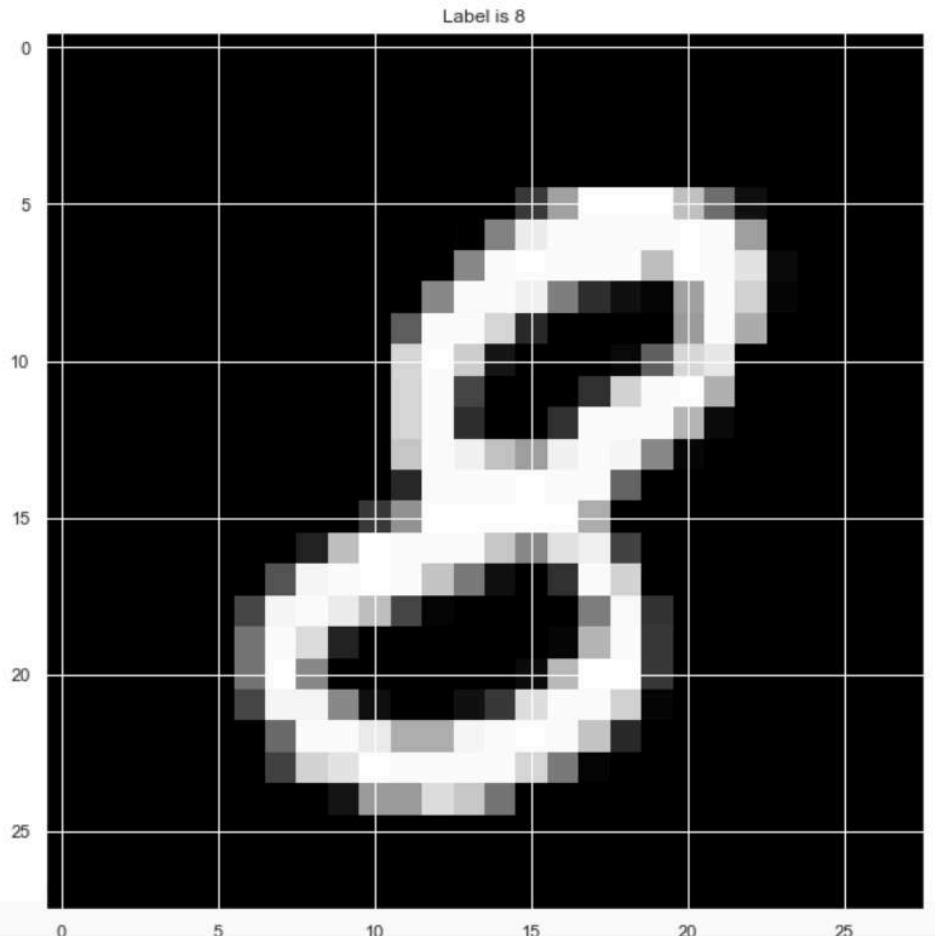
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)

# Visualization Candies
import matplotlib.pyplot as plt

def show_image(mnist_x_vector, mnist_y_vector):
    pixels = mnist_x_vector.reshape((28, 28))
    label = np.where(mnist_y_vector == 1)[0][0]
    plt.title('Label is {}'.format(label))
    plt.imshow(pixels, cmap='gray')
    plt.show()

# Fifth image and label.
show_image(mnist.train.images[5], mnist.train.labels[5])

```



```
X_mnist = torch.tensor(mnist.train.images).float()
Y_mnist = torch.tensor(mnist.train.labels).float()

X_mnist_test = torch.tensor(mnist.test.images).float()
Y_mnist_test = torch.tensor(mnist.test.labels).float()
```

Fill in the blanks =)

```
# Use FloatTensor.shape to get the shape of the matrix/tensor.
num_data, input_dim = X_mnist.shape
print('Inputs Dim:', input_dim)

num_data, output_dim = ???
print('Output Dim:', output_dim)
```

```
Inputs Dim: 784
Output Dim: 10
```

```
hidden_dim = 500

model = nn.Sequential(nn.Linear(???, ???),
                      nn.???)()

criterion = nn.???
learning_rate = 1.0
optimizer = optim.???(model.parameters(), lr=learning_rate)
num_epochs = 100

losses = []
plt.ion()

for _e in tqdm(range(num_epochs)):
    optimizer.zero_grad()
    predictions = ???
    loss_this_epoch = ???
    loss_this_epoch.???
    optimizer.???
    ##print([float(_pred) for _pred in predictions], list(map(int, Y_pt)), loss_this_epoch.data[0])
    losses.append(loss_this_epoch.data.item())

    clear_output(wait=True)
    plt.plot(losses)
    plt.pause(0.05)
```

```
X_mnist = torch.tensor(mnist.train.images).float()
Y_mnist = torch.tensor(mnist.train.labels).float()

X_mnist_test = torch.tensor(mnist.test.images).float()
Y_mnist_test = torch.tensor(mnist.test.labels).float()
```

```
# Use FloatTensor.shape to get the shape of the matrix/tensor.
num_data, input_dim = X_mnist.shape
print('Inputs Dim:', input_dim)

num_data, output_dim = Y_mnist.shape
print('Output Dim:', output_dim)
```

```
Inputs Dim: 784
Output Dim: 10
```

```
hidden_dim = 500

model = nn.Sequential(nn.Linear(input_dim, output_dim),
                      nn.Sigmoid())

criterion = nn.MSELoss()
learning_rate = 1.0
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
num_epochs = 100

losses = []
plt.ion()

for _e in tqdm(range(num_epochs)):
    optimizer.zero_grad()
    predictions = model(X_mnist)
    loss_this_epoch = criterion(predictions, Y_mnist)
    loss_this_epoch.backward()
    optimizer.step()
    ##print([float(_pred) for _pred in predictions], list(map(int, Y_pt)), loss_this_epoch.data[0])
    losses.append(loss_this_epoch.data.item())

    clear_output(wait=True)
    plt.plot(losses)
    plt.pause(0.05)
```

```
X_mnist = torch.tensor(mnist.train.images).float()
Y_mnist = torch.tensor(mnist.train.labels).float()

X_mnist_test = torch.tensor(mnist.test.images).float()
Y_mnist_test = torch.tensor(mnist.test.labels).float()
```

```
# Use FloatTensor.shape to get the shape of the matrix/tensor.
num_data, input_dim = X_mnist.shape
print('Inputs Dim:', input_dim)

num_data, output_dim = Y_mnist.shape
print('Output Dim:', output_dim)
```

```
Inputs Dim: 784
Output Dim: 10
```

```
hidden_dim = 500

model = nn.Sequential(nn.Linear(input_dim, output_dim),
                      nn.Sigmoid())

criterion = nn.MSELoss()
learning_rate = 1.0
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
num_epochs = 100

losses = []
plt.ion()

for _e in tqdm(range(num_epochs)):
    optimizer.zero_grad()
    predictions = model(X_mnist)
    loss_this_epoch = criterion(predictions, Y_mnist)
    loss_this_epoch.backward()
    optimizer.step()
    ##print([float(_pred) for _pred in predictions], list(map(int, Y_pt)), loss_this_epoch.data[0])
    losses.append(loss_this_epoch.data.item())

    clear_output(wait=True)
    plt.plot(losses)
    plt.pause(0.05)
```

```
predictions = model(X_mnist_test)
predictions

tensor([[0.1220, 0.0558, 0.1120, ..., 0.8124, 0.1212, 0.2739],
       [0.1745, 0.1066, 0.4185, ..., 0.0175, 0.1121, 0.0340],
       [0.0963, 0.7832, 0.2290, ..., 0.1862, 0.1943, 0.1618],
       ...,
       [0.0073, 0.0296, 0.0250, ..., 0.1103, 0.0635, 0.1627],
       [0.0748, 0.1572, 0.0525, ..., 0.0713, 0.1085, 0.0661],
       [0.1785, 0.0087, 0.1558, ..., 0.0084, 0.0274, 0.0240]],
      grad_fn=<SigmoidBackward>)

pred = np.array([np.argmax(_p) for _p in predictions.data.numpy()])
pred

array([7, 2, 1, ..., 4, 1, 6])

truth = np.array([np.argmax(_p) for _p in Y_mnist_test.data.numpy()])
truth

array([7, 2, 1, ..., 4, 5, 6])

(pred == truth).sum() / len(pred)
```

0.7895

```
hidden_dim = 500

model = nn.Sequential(nn.Linear(input_dim, hidden_dim),
                      nn.Sigmoid(),
                      nn.Linear(hidden_dim, output_dim),
                      nn.Sigmoid())

criterion = nn.MSELoss()
learning_rate = 1.0
optimizer = optim.SGD(model.parameters(), lr=learning_rate)
num_epochs = 1000

losses = []
plt.ion()

for _e in tqdm(range(num_epochs)):
    optimizer.zero_grad()
    predictions = model(X_mnist)
    loss_this_epoch = criterion(predictions, Y_mnist)
    loss_this_epoch.backward()
    optimizer.step()
    ##print([float(_pred) for _pred in predictions], list(map(int, Y_pt)), loss_this_epoch.data[0])
    losses.append(loss_this_epoch.data.item())

plt.plot(losses)
```

2% | 21/1000 [00:35<27:24, 1.68s/it]