



Text Processing using Machine Learning

More Deep Learning Foundations

Liling Tan

02-03 Dec 2019

OVER **5,500** GRADUATE ALUMNI OFFERING OVER **120** ENTERPRISE IT, INNOVATION & LEADERSHIP PROGRAMMES TRAINING OVER **120,000** DIGITAL LEADERS & PROFESSIONALS

Lecture

- **More Deep Learning Foundations**
 - Recap: Behind the PyTorch Magic
 - Activation Functions
 - Hands-on

More Deep Learning Foundations

Recap and activation functions

Recap: XOR with PyTorch

```
2 import torch
3 from torch import nn
4
5 X = xor_input = torch.tensor([[0,0], [0,1], [1,0], [1,1]]).float().to(device)
6 Y = xor_output = torch.tensor([[0], [1], [1], [0]]).float().to(device)
7
8 # Define the shape of the weight vector.
9 num_data, input_dim = X.shape
10 hidden_dim = 5
11 output_dim = len(Y)
12
13 # Initialize the network as an nn.Sequential.
14 model = nn.Sequential(
15     nn.Linear(input_dim, hidden_dim),
16     nn.Sigmoid(),
17     nn.Linear(hidden_dim, output_dim),
18     nn.Sigmoid()
19 ).to(device)
20
21 # Declare the loss function as an nn.Module.
22 criterion = nn.MSELoss()
```

Recap: XOR with PyTorch

```
24 num_epochs = 10000
25 learning_rate = 0.3
26
27 # Initialize the optimizer that'll help us with the parameters updates.
28 optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
29
30 losses = []
31 for epoch_n in tqdm(range(num_epochs)):
32     # When we move on in this epoch, we should forget the .grad tensors
33     # in the previous epoch since we already made the updates.
34     optimizer.zero_grad()
35
36     # Forward propagation.
37     predictions = model(X)
38     # Compute Loss
39     loss = criterion(predictions, Y)
40     # Keep track of the losses.
41     losses.append(loss.item())
42
43     # Backpropagation.
44     loss.backward()
45
46     # The step() function will update the parameters in the models that
47     # has the .grad tensors respectively.
48     optimizer.step()
```

**Forward
propagation
magic!!!**



Recap: XOR from Scratch

```

3  def sigmoid(x): # Returns values that sums to one.
4      return 1 / (1 + np.exp(-x))
5
6  X = xor_input = np.array([[0,0], [0,1], [1,0], [1,1]])
7  Y = xor_output = np.array([[0,1,1,0]]).T
8
9  # Define the shape of the weight vector.
10 num_data, input_dim = X.shape
11 hidden_dim = 5
12 output_dim = len(Y.T)
13
14 # Initialize weights between the input layers and the hidden layer.
15 W1 = np.random.random((input_dim, hidden_dim))
16 # Initialize weights between the hidden layers and the output layer.
17 W2 = np.random.random((hidden_dim, output_dim))
18
19 # Initialize weigh
20 num_epochs = 5000
21 learning_rate = 0.15
22
23 for epoch_n in range(num_epochs):
24     layer0 = X
25     # Inside the perceptron, Step 2.
26     layer1 = sigmoid(np.dot(layer0, W1))
27     layer2 = sigmoid(np.dot(layer1, W2))
  
```

**Manually
doing matrix
multiplication**



Recap: XOR with PyTorch (BTS)

```

2  def sigmoid(x): # Returns values that sums to one.
3      return 1 / (1 + torch.exp(-x))
4
5  X = xor_input = torch.tensor([[0,0], [0,1], [1,0], [1,1]]).float().to(device)
6  Y = xor_output = torch.tensor([[0], [1], [1], [0]]).float().to(device)
7
8  # Define the shape of the weight vector.
9  num_data, input_dim = X.shape
10 hidden_dim = 5
11 output_dim = len(Y)
12 # When we initialize tensors that needs updating, we use `requires_grad=True`
13 # for autograd to kick in later on.
14 W1 = torch.randn(input_dim, hidden_dim, requires_grad=True).to(device)
15 W2 = torch.randn(hidden_dim, output_dim, requires_grad=True).to(device)
16
17 num_epochs = 10000
18 learning_rate = 0.3
19
20 for epoch_n in tqdm(range(num_epochs)):
21     layer0 = X
22     # See https://pytorch.org/docs/stable/torch.html#torch-mm
23     # Use the torch.tensor.mm() instead of np.dot()
24     layer1 = sigmoid(X.mm(W1))
25     layer2 = sigmoid(layer1.mm(W2))
  
```

**PyTorch is
actually
doing matrix
multiplication
Behind-The-
Scene (BTS)**

Recap: XOR with PyTorch

```
24 num_epochs = 10000
25 learning_rate = 0.3
26
27 # Initialize the optimizer that'll help us with the parameters updates.
28 optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
29
30 losses = []
31 for epoch_n in tqdm(range(num_epochs)):
32     # When we move on in this epoch, we should forget the .grad tensors
33     # in the previous epoch since we already made the updates.
34     optimizer.zero_grad()
35
36     # Forward propagation.
37     predictions = model(X)
38     # Compute Loss
39     loss = criterion(predictions, Y)
40     # Keep track of the losses.
41     losses.append(loss.item())
42
43     # Backpropagation.
44     loss.backward()
45
46     # The step() function will update the parameters in the models that
47     # has the .grad tensors respectively.
48     optimizer.step()
```

**Backprop
magic!!!**



Recap: XOR from Scratch

```
for epoch_n in range(num_epochs):
    layer0 = X
    # Forward propagation.

    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(layer0, W1))
    layer2 = sigmoid(np.dot(layer1, W2))

    # Back propagation (Y -> layer2)
    # How much did we miss in the predictions?
    cost_error = mse(layer2, Y)

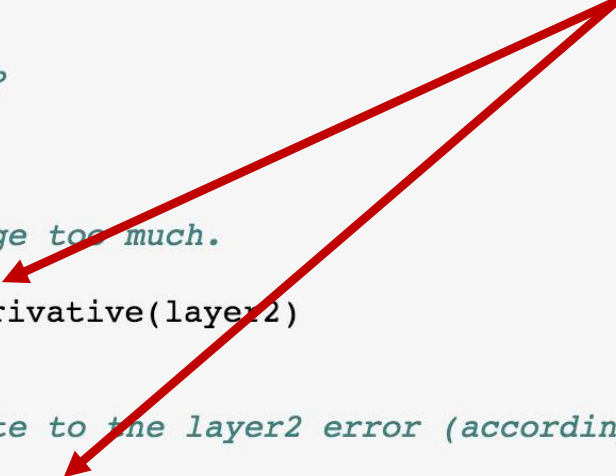
    # In what direction is the target value?
    # Were we really close? If so, don't change too much.
    layer2_error = mse_derivative(layer2, Y)
    layer2_delta = layer2_error * sigmoid_derivative(layer2)

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer2 error (according to the weights)?
    layer1_error = np.dot(layer2_delta, W2.T)
    layer1_delta = layer1_error * sigmoid_derivative(layer1)


    # update weights
    W2 += - learning_rate * np.dot(layer1.T, layer2_delta)
    W1 += - learning_rate * np.dot(layer0.T, layer1_delta)
    #print(np.dot(layer0.T, layer1_delta))
    #print(epoch_n, list((layer2)))

    # Log the loss value as we proceed through the epochs.
    losses.append(cost_error)
    #print(cost_delta)
```

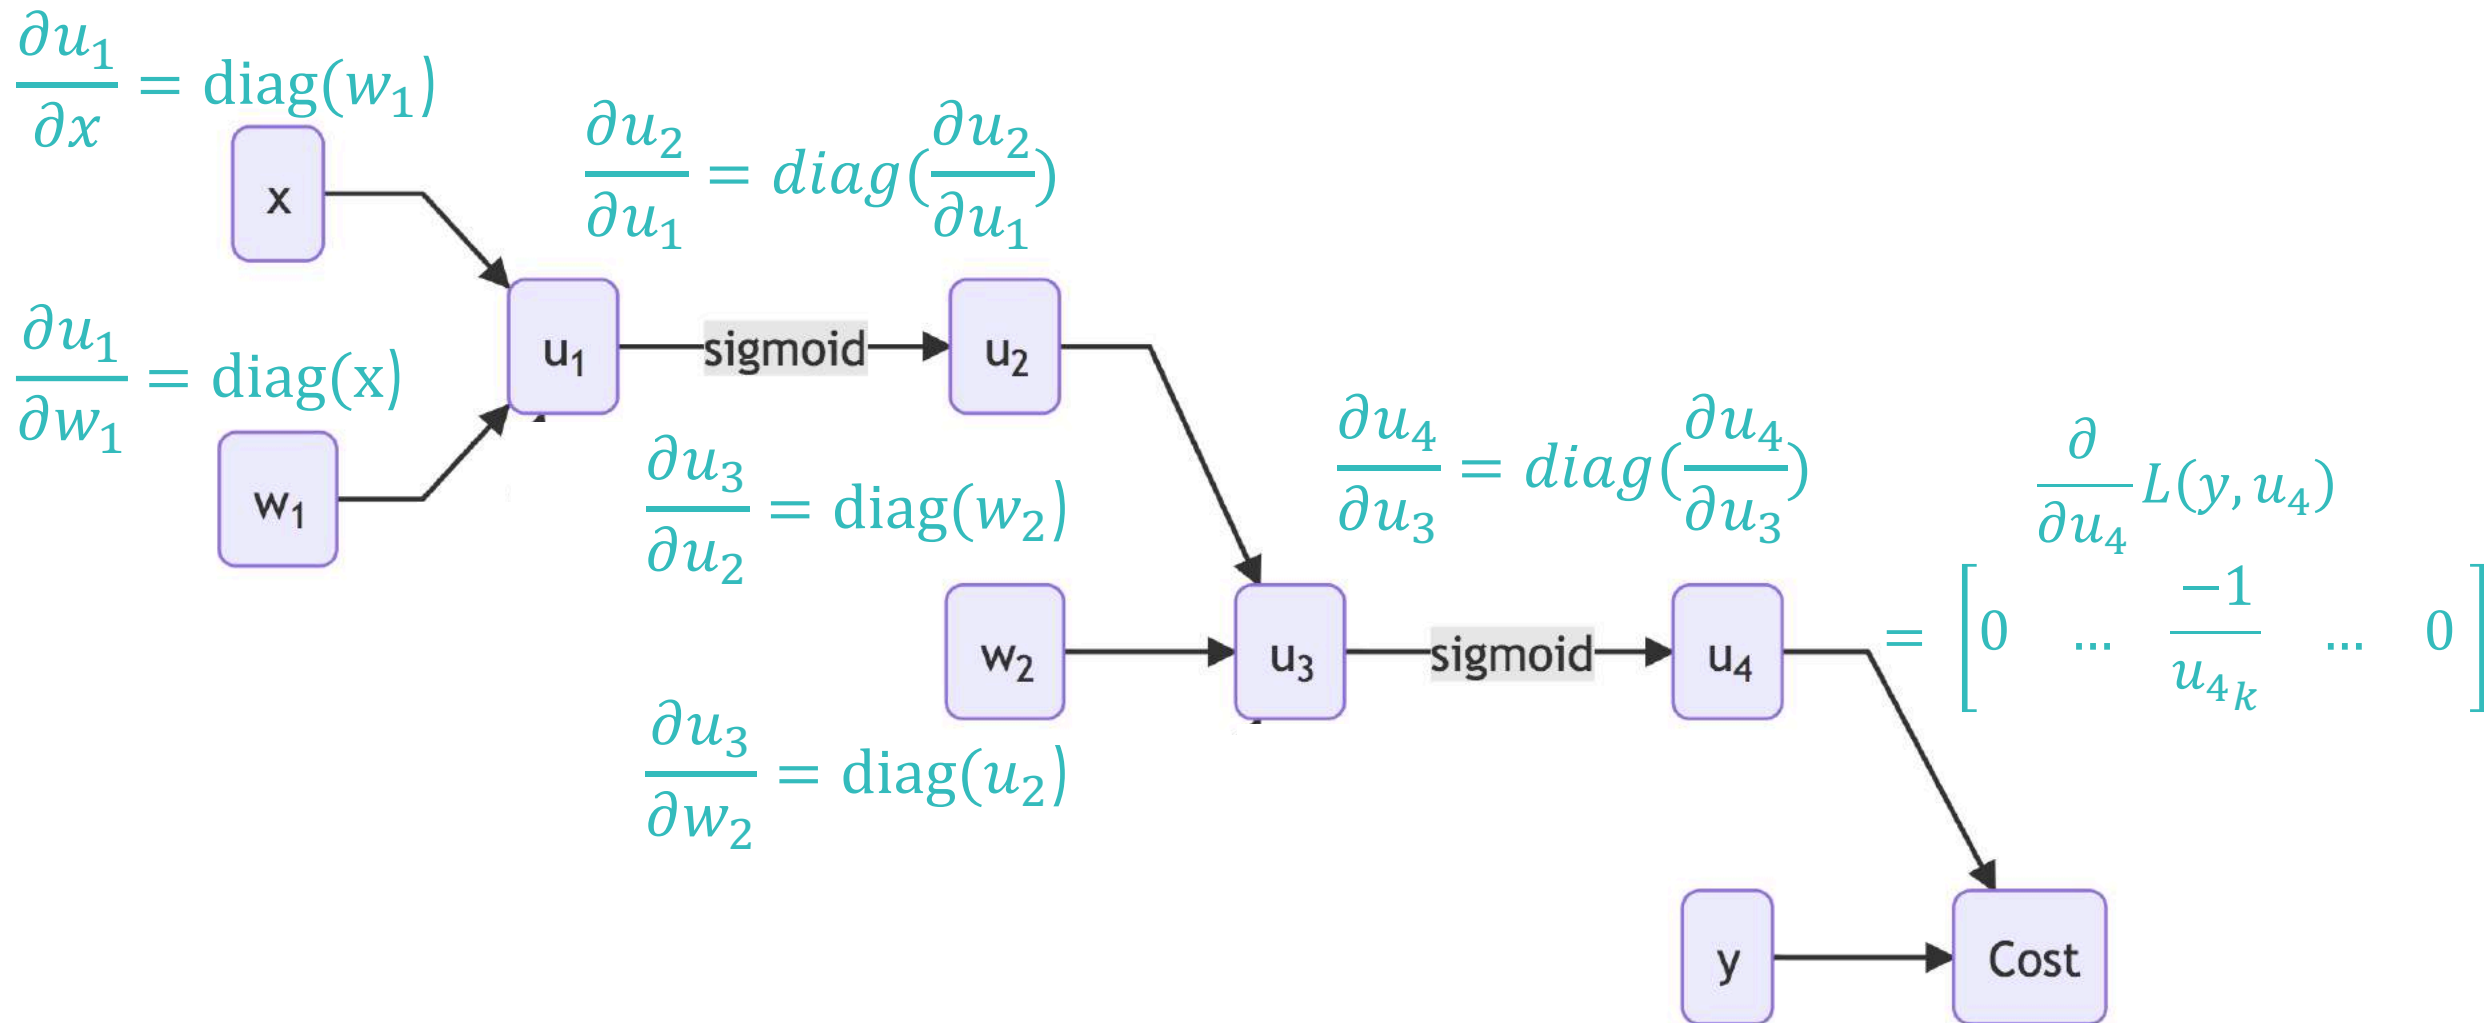
Backprop
pain!!!



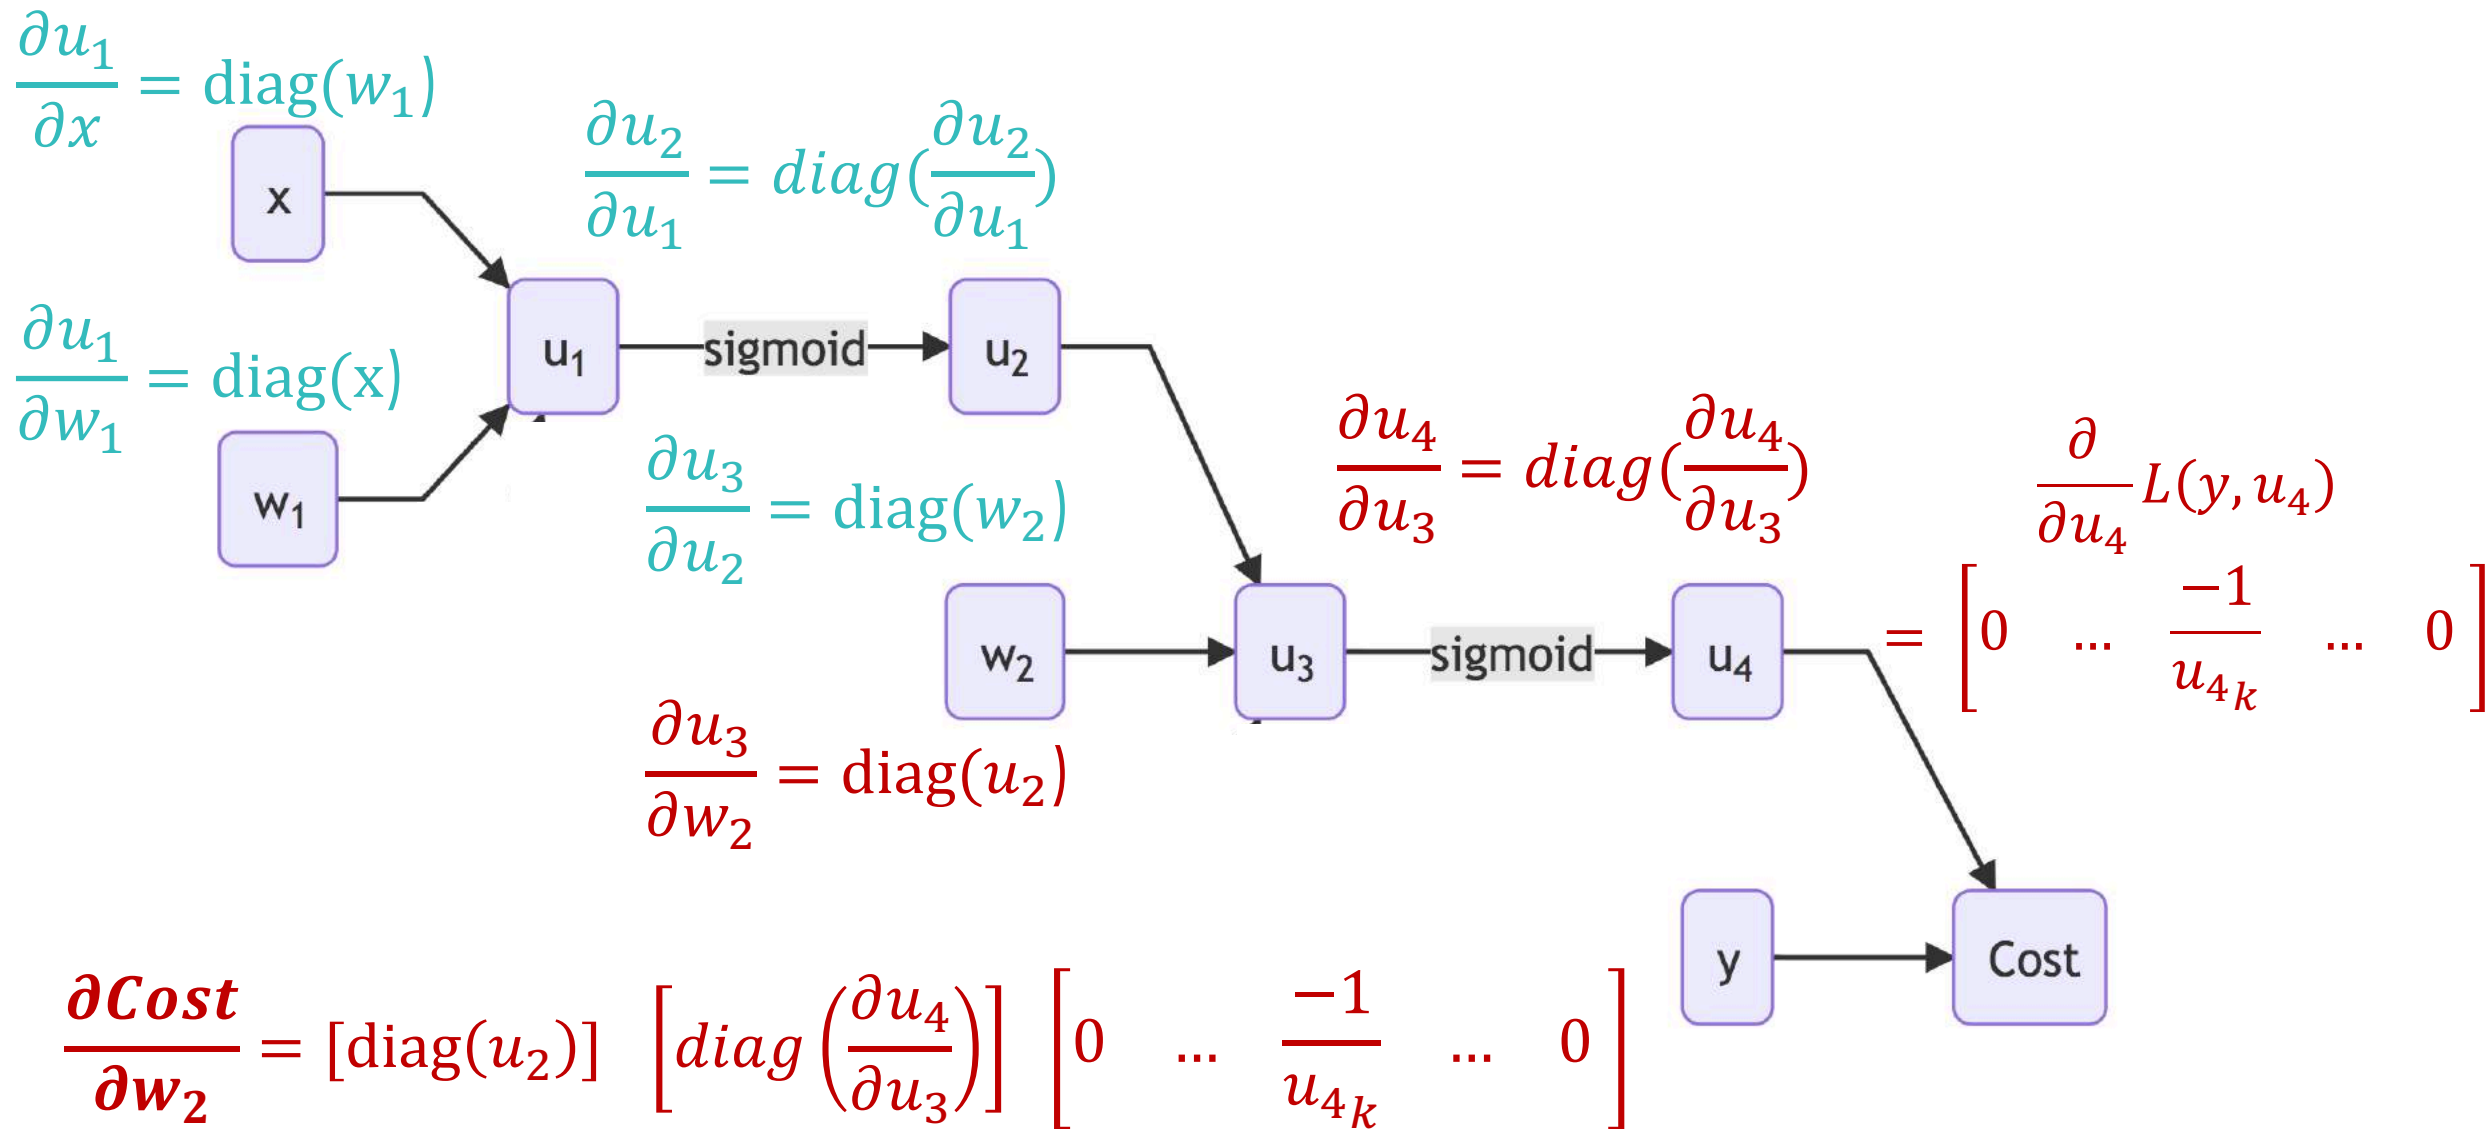
Recap: Derivatives of a Multi-Layered Perceptron



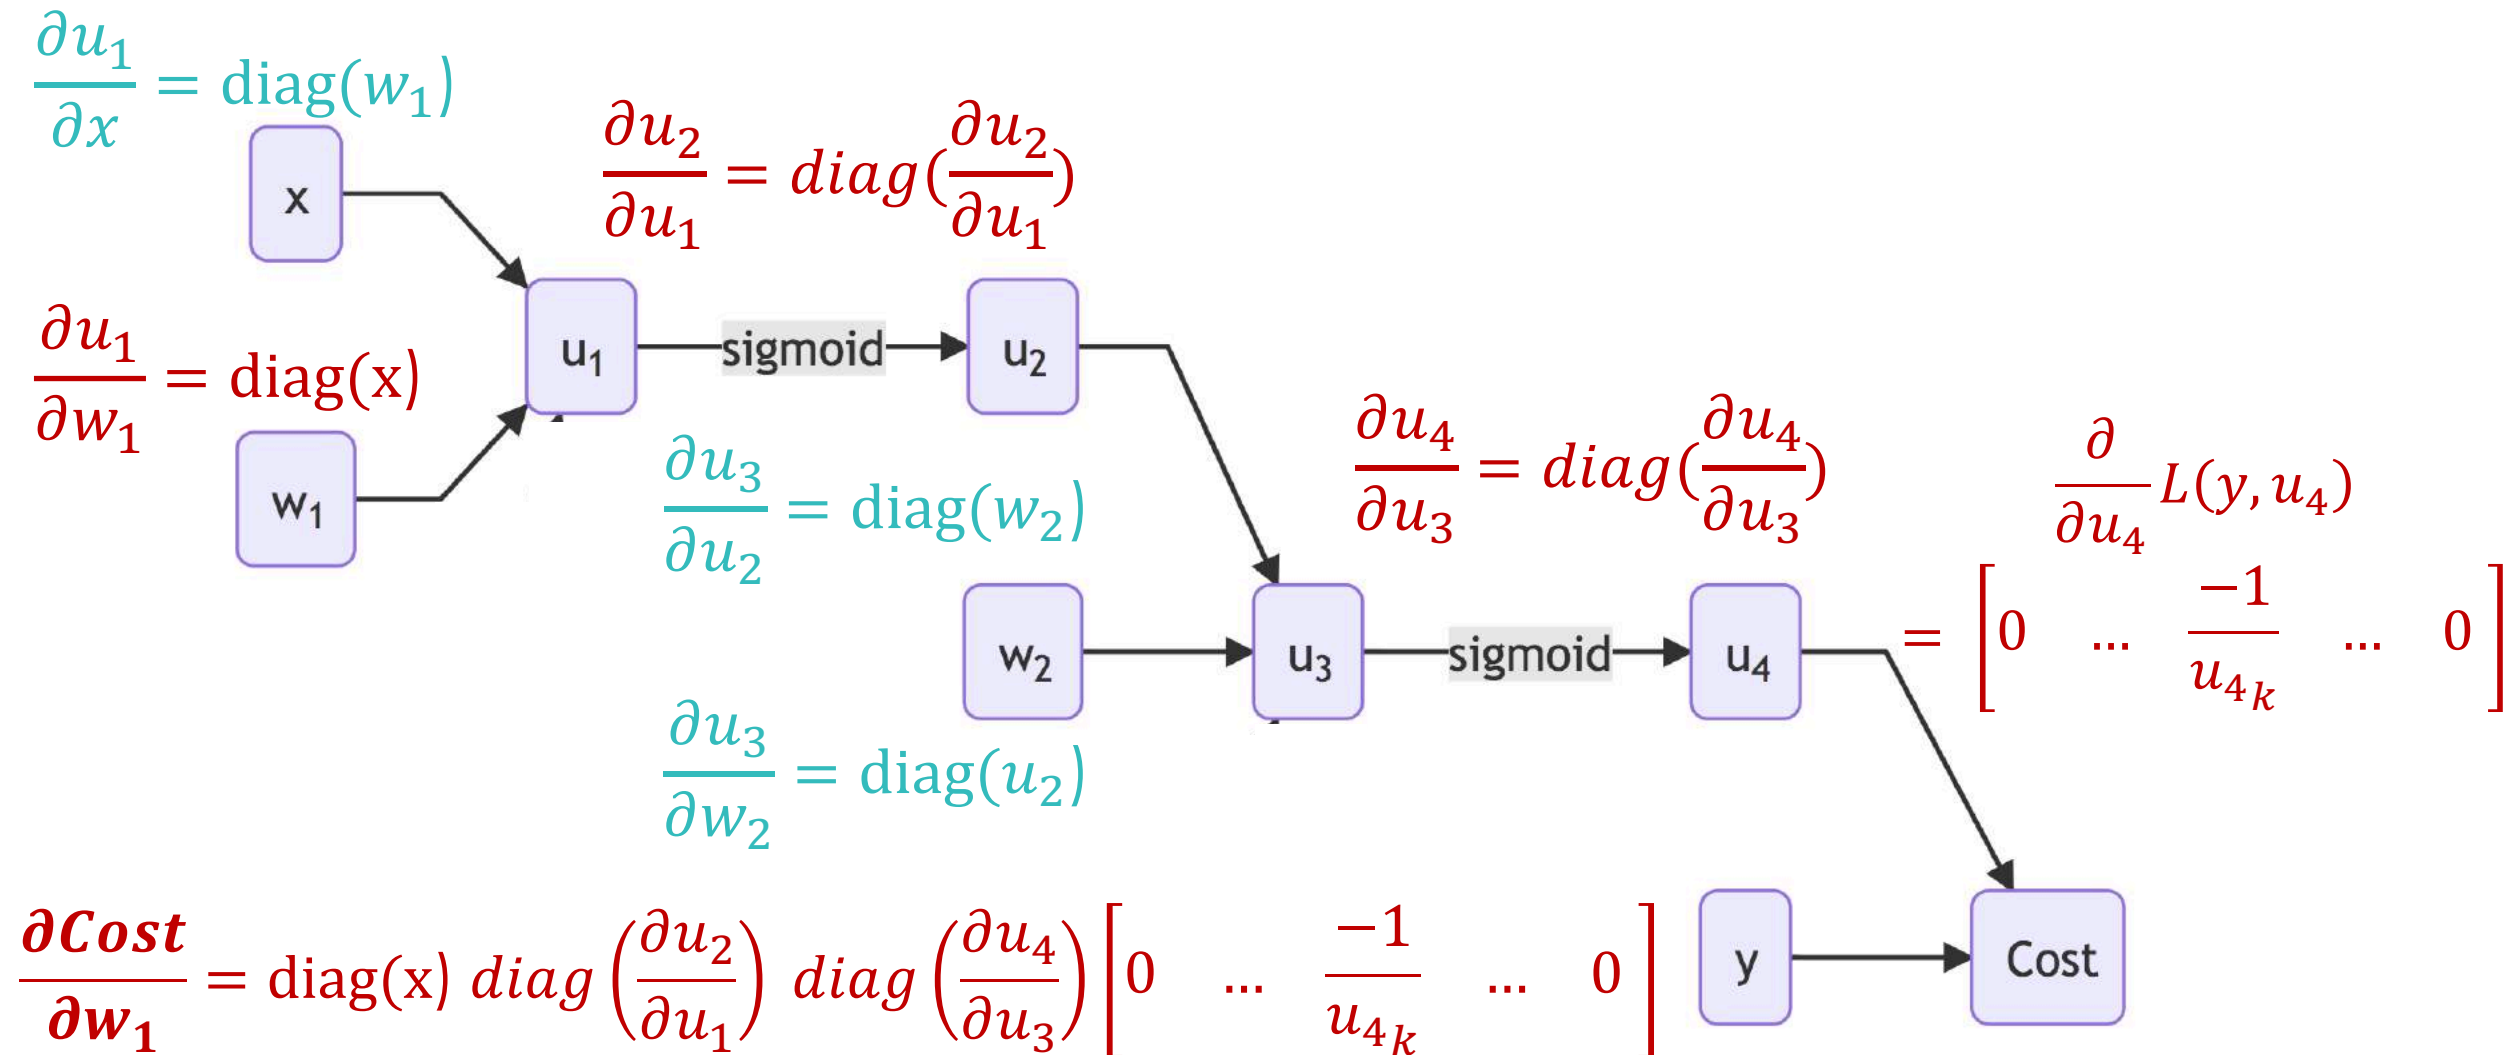
Recap: Derivatives of a Multi-Layered Perceptron



Recap: Derivatives of a Multi-Layered Perceptron



Recap: Derivatives of a Multi-Layered Perceptron



Recap: Derivatives of a Multi-Layered Perceptron

$$\frac{\partial Cost}{\partial w_1} = \text{diag}(x) \text{diag}\left(\frac{\partial u_2}{\partial u_1}\right) \text{diag}\left(\frac{\partial u_4}{\partial u_3}\right) \begin{bmatrix} 0 & \dots & \frac{-1}{u_{4k}} & \dots & 0 \end{bmatrix}$$

$$\frac{\partial Cost}{\partial w_2} = \text{diag}(u_2) \text{diag}\left(\frac{\partial u_4}{\partial u_3}\right) \begin{bmatrix} 0 & \dots & \frac{-1}{u_{4k}} & \dots & 0 \end{bmatrix}$$

$$w_2 += -lr * \frac{\partial Cost}{\partial w_2}$$

$$w_1 += -lr * \frac{\partial Cost}{\partial w_1}$$

Recap: XOR with PyTorch

```
24 num_epochs = 10000
25 learning_rate = 0.3
26
27 # Initialize the optimizer that'll help us with the parameters updates.
28 optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
29
30 losses = []
31 for epoch_n in tqdm(range(num_epochs)):
32     # When we move on in this epoch, we should forget the .grad tensors
33     # in the previous epoch since we already made the updates.
34     optimizer.zero_grad()
35
36     # Forward propagation.
37     predictions = model(X)
38     # Compute Loss
39     loss = criterion(predictions, Y)
40     # Keep track of the losses.
41     losses.append(loss.item())
42
43     # Backpropagation.
44     loss.backward()
45
46     # The step() function will update the parameters in the models that
47     # has the .grad tensors respectively.
48     optimizer.step()
```

Optimizer
weights
updating
magic!!!

Recap: XOR from Scratch

```
for epoch_n in range(num_epochs):
    layer0 = X
    # Forward propagation.

    # Inside the perceptron, Step 2.
    layer1 = sigmoid(np.dot(layer0, W1))
    layer2 = sigmoid(np.dot(layer1, W2))

    # Back propagation (Y -> layer2)
    # How much did we miss in the predictions?
    cost_error = mse(layer2, Y)

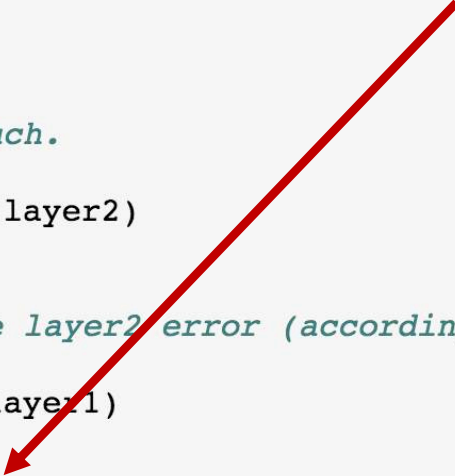
    # In what direction is the target value?
    # Were we really close? If so, don't change too much.
    layer2_error = mse_derivative(layer2, Y)
    layer2_delta = layer2_error * sigmoid_derivative(layer2)

    # Back propagation (layer2 -> layer1)
    # How much did each layer1 value contribute to the layer2 error (according to the weights)?
    layer1_error = np.dot(layer2_delta, W2.T)
    layer1_delta = layer1_error * sigmoid_derivative(layer1)

    # update weights
    W2 += - learning_rate * np.dot(layer1.T, layer2_delta)
    W1 += - learning_rate * np.dot(layer0.T, layer1_delta)
    #print(np.dot(layer0.T, layer1_delta))
    #print(epoch_n, list((layer2)))

    # Log the loss value as we proceed through the epochs.
    losses.append(cost_error)
    #print(cost_delta)
```

Painful
weights
updates



Recap: XOR with PyTorch (BTS)

```
2  for epoch_n in tqdm(range(num_epochs)):
3      layer0 = X
4      layer1 = sigmoid(X.mm(W1))
5      layer2 = sigmoid(layer1.mm(W2))
6
7      # Loss is a Tensor of torch.Size([]), so and loss.item() is a scalar/float.
8      # Try printing `print(loss.shape)` to confirm the above.
9      loss = mse(layer2, Y)
10     # Keep track of the losses.
11     losses.append(loss.item())
12
13     # The `loss.backward()` will compute the gradient of loss w.r.t. all
14     # tensors that has `requires_grad=True`.
15     # After this, W1.grad and W2.grad will hold the gradients of
16     # the loss w.r.t. to W1 and W2 respectively.
17     loss.backward()
18
19     # Now we have the backpropagated gradients, we want to update the weights.
20     # Whenever you perform tensor operations on tensors that has `requires_grad=True`,
21     # pytorch will try to build computation graph. For now, we only need to
22     # force the updates of our weights without forming more computation graph,
23     # so we use the no_grad() context manager:
24     with torch.no_grad():
25         W1 += -learning_rate * W1.grad
26         W2 += -learning_rate * W2.grad
27
```


**PyTorch
Optimizer
weights
updating
Behind-The-
Scene (BTS)**

Activation Functions

Sigmoid, S-Shape, Rectified Activations

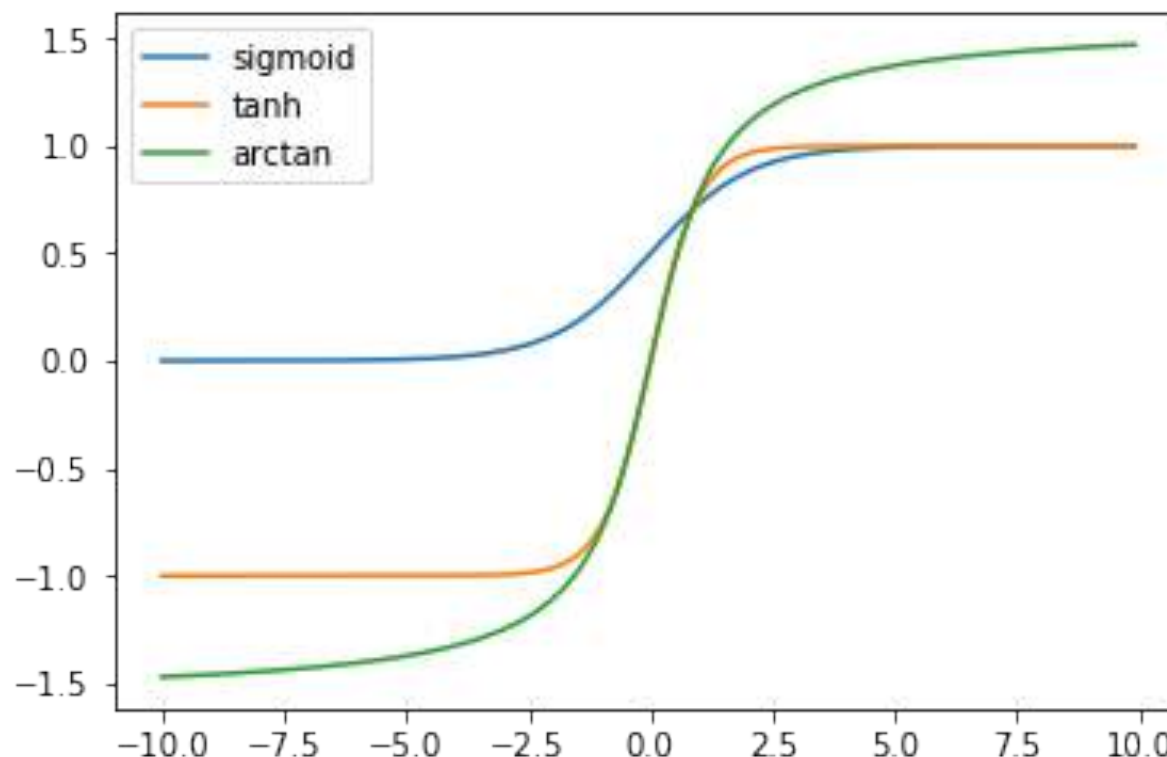
Activation Function (Sigmoid)

```
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 def sigmoid(x):
7     return 1/(1+np.exp(-x))
8
9 # Generate points from -10 to +10,
10 # in steps of 0.1
11 x = np.arange(-10, 10, 0.1)
12 y = sigmoid(x)
13
14 # Plot the graph.
15 plt.plot(x, y)
16 plt.show()
```




Activation Function (S-Shape Activations)

```
6 def sigmoid(x):
7     return 1 / (1+np.exp(-x))
8
9 def tanh(x):
10    return np.tanh(x)
11
12 def arctan(x):
13    return np.arctan(x)
14
15 x = np.arange(-10, 10, 0.1)
16 y1 = sigmoid(x)
17 y2 = tanh(x)
18 y3 = arctan(x)
19
20
21 plt.plot(x,y1, label='sigmoid')
22 plt.plot(x,y2, label='tanh')
23 plt.plot(x,y3, label='arctan')
24 plt.legend(loc='upper left')
25 plt.show()
```



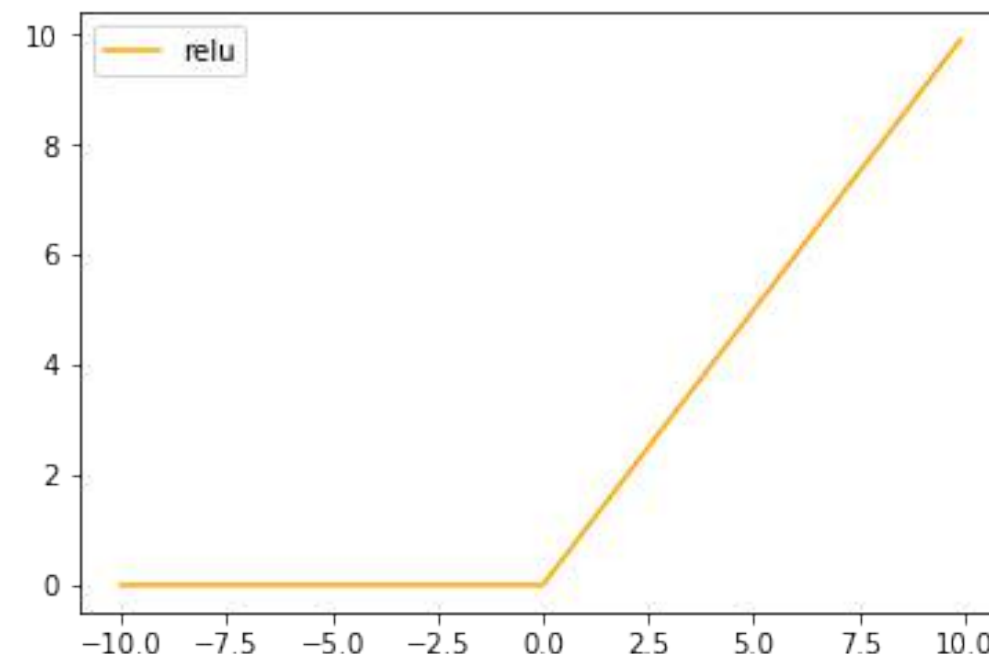
Activation Function (Softmax)

```
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 def softmax(x):
7     return np.exp(x) / np.sum(np.exp(x), axis=0)
8
9 x = np.arange(-10, 10, 0.1)
10 y = softmax(x)
11
12 plt.plot(x,y)
13 plt.show()
```




Activation Function (ReLU)

```
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 def relu(x):
7     return x * (x > 0)
8
9 y2 = relu(x)
10
11 plt.plot(x,y2, label='relu', color='orange')
12 plt.legend(loc='upper left')
13 plt.show()
```



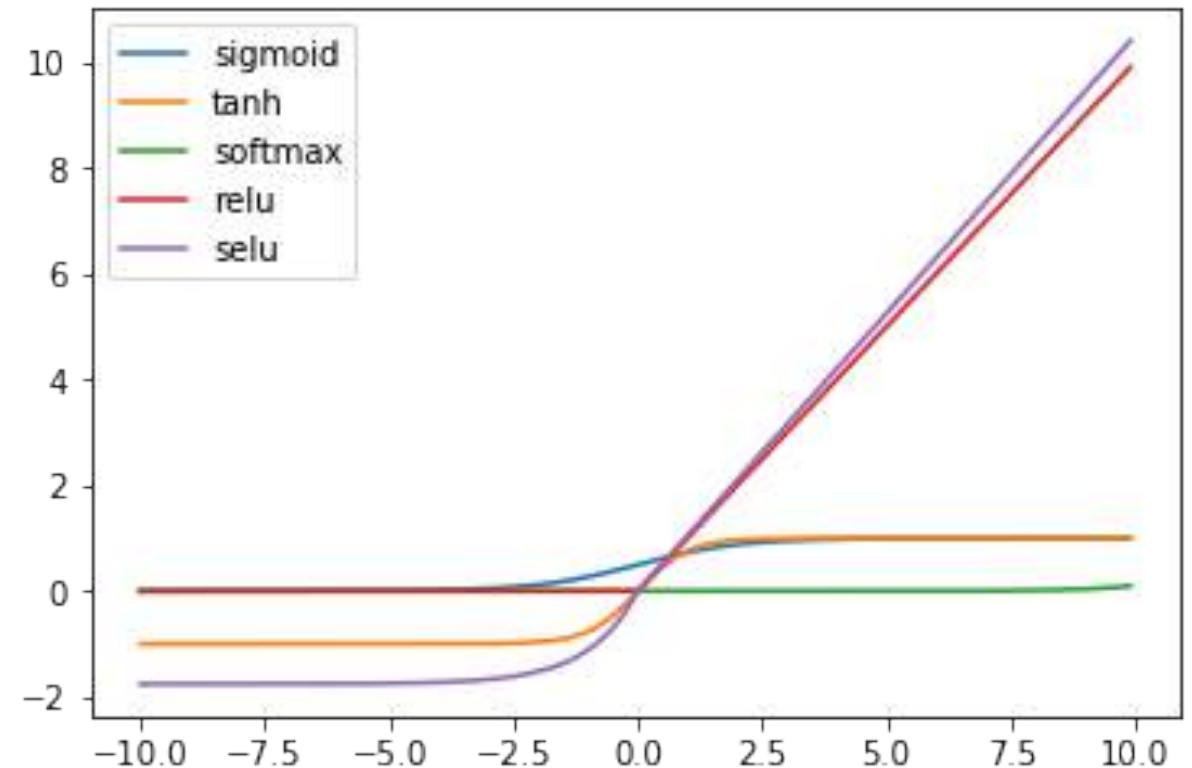
Activation Functions (SELU)

```
3 def selu(x):
4     alpha = 1.6732632423543772848170429916717
5     scale = 1.0507009873554804934193349852946
6     return (np.maximum(0, x) +
7             np.minimum(0, alpha*(np.exp(x) -1)))
8
9 y3 = selu(x)
10
11 plt.plot(x,y3, label='selu', color='green')
12 plt.legend(loc='upper left')
13 plt.show()
```



Activation Functions with PyTorch

```
3  from torch import nn, tensor
4
5  x = tensor(np.arange(-10, 10, 0.1))
6
7  a1 = nn.Sigmoid()
8  a2 = nn.Tanh()
9  a3 = nn.Softmax()
10 a4 = nn.ReLU()
11 a5 = nn.SELU()
12
13 y1, y2, y3 = a1(x), a2(x), a3(x)
14 y4, y5 = a4(x), a5(x)
15
16 plt.plot(x, y1, label='sigmoid')
17 plt.plot(x, y2, label='tanh')
18 plt.plot(x, y3, label='softmax')
19 plt.plot(x, y4, label='relu')
20 plt.plot(x, y5, label='selu')
21 plt.legend(loc='upper left')
22 plt.show()
```



Hands-on: Unmagical PyTorch

Possibly XOR one more time =)

Environment Setup

Open Anaconda Navigator.

Go to the PyTorch installation page, copy the command as per configuration:

<https://pytorch.org/get-started/locally/>

Fire up the terminal in Anaconda Navigator.

Start a Jupyter Notebook.

Download <http://bit.ly/ANLP-Session2-Empty->

Import the .ipynb to the Jupyter Notebook

Fín