ISO 29990:2010

NATIONAL INFOCOMM COMPETENCY FRAMEWORK

NUS National University of Singapore | ISS INSTITUTE OF SYSTEMS SCIENCE

Text Processing using Machine Learning

Ask Me Anything

Liling Tan

2019

OVER 5,500 GRADUATE ALUMNI

OFFERING OVER 120 ENTERPRISE IT, INNOVATION & LEADERSHIP PROGRAMMES

TRAINING OVER 120,000 DIGITAL LEADERS & PROFESSIONALS

© 2018 National University of Singapore. All Rights Reserved

Ask Me Anything

© 2018 National University of Singapore. All Rights Reserved

Frequently Asked Questions

- What's next in Machine Learning and NLP?
- Where do we learn more on ML/DL and NLP?
- How do I get industrial experience for ML/DL and NLP?
- What's the latest thing in NLP now?
- Tell me about your daily work, how does a day as "research scientist" look like?
- I have problem X, which tool/method Y to solve it?
- Open Source Software

© 2018 National University of Singapore. All Rights Reserved

What's next in Machine Learning and NLP (that I'm personally excited about)?

- **Quantum Computing**
  - https://medium.com/xanaduai/training-quantum-neural-networks-with-pennylane-pytorch-and-tensorflow-c669108118cc
  - https://github.com/XanaduAI/pennylane
- **Federated Learning**
  - https://blog.openmined.org/upgrade-to-federated-learning-in-10-lines
  - https://github.com/OpenMined/PySyft
- **Evolutionary/Genetic Algorithms**
  - https://leanpub.com/genetic_algorithms_with_python

© 2018 National University of Singapore. All Rights Reserved

Where do we learn more on ML/DL and NLP?

- **As a start**
  - https://github.com/datasciencesg/workshops/tree/master/LearnItYourself

- **For latest shiny things**
  - Follow DL/ML stars on Twitter or their preferred social media
  - Follow the #nlproc hashtag on Twitter
  - Follow @arxiv_cs_cl on Twitter and for sanity http://www.arxiv-sanity.com/

- **Publish a paper and/or Join a conference**
  - Most listed conferences on https://aclanthology.info/ are good to join
  - Text, Speech, and Dialogue (TSD) and Interspeech conferences
  - Join shared task in workshops co-located in these conferences

© 2018 National University of Singapore. All Rights Reserved

How to get industrial experience for ML/DL & NLP?

- **Join Competitions**
  - Kaggle, shared task in conferences and many more
- **Build things and open source**
  - Learn some Flask/Django or web development, just enough to show the world and demo what you've done.
- **Get a mentor or an internship**
  - Mentorship is harder to find but it's possible. Sometimes non-profit organizations and companies do have mentorship programs
  - Internships are a plenty but find places that don't make you do "sia kang" and people you think would enjoy working with

© 2018 National University of Singapore. All Rights Reserved

What's the latest thing in NLP now?

- **Transformers**
  - Lots of transformers and its variant
  - I do want to see it go away… It's sort of a boring model.

- **You'll never know, it moves so fast…**
  - Every day new code commits are made on PyTorch, Tensorflow, AllenNLP, SpaCy, etc.
  - Get involve in the open source and you get first blood on the new tech =)
  - Every 2-3 months is an NLP conference, every month there's an NLP conference deadline

- **Don't chase the shiny new things**
  - Know that they exist, know how they work and what libraries to use
  - Use it only when you tried and experimented and show that it works better for your task
  - Know the foundations, there's seldom something new under the sun, just better rehashes of things

© 2018 National University of Singapore. All Rights Reserved

Research Scientist @ RIT

- **Project Management**
  - Managing expectations of "AI" products
  - Understanding what problem your "clients" wants to solve
  - Know what data they have/have not, find where to get the required data
  - Propose a feasible solution and try before 2nd meeting
  - Keep "clients" engaged, show to them it's the latest tech that's useful to them

- **Knowing Backend/Frontend Engineering helps**
  - Know what's possible, what's easy what's hard
  - Learn from engineers (dockers, databases, cloud, apps design/dev) and let engineers learn what you do
  - There's not clear boundary, a data/dev engineer might train a better model than you do

- **Reading and lots of coding**
  - Code sprints to get **** done
  - Finding out what's new, useful and quick to prototype
  - Knowing whether it can be "productionize", e.g. ensemble of 100+ models wins competition -_-| | |
  - Look at old ideas, know the limitations, see how you can fix them

I have problem X, which tool/method Y to solve it?

- **Literature review**
- **Know what are the datasets and what's in them (noise, quirks, etc.)**
- **Which evaluation metric is task X evaluated on?**
- **Find the latest shiniest paper,**
  - Track the oldest relevant citation of the task, read that paper
  - Find the highest cited paper for the task, use that as your baseline
  - Whenever possible, hunt down the datasets in that highest cited paper and latest shiniest paper
  - Define your success criteria for the task industrially (it might not be the standard eval metric for the task)
  - Try/Reimplement the baseline
  - Did baseline meet the success criteria? Can your engineer productionize it?
  - Ask the business/project stakeholder whether it's sufficient
  - Communicate your model/libraries to engineers, build it, test it, break it, repeat

Open Source Software

- *"If data is the fuel to today's software, open source is the fire."*
  – Nat Gillin

- **Learn and learning a lot from reading code, fixing bugs, testing things, reimplementing stuff**

- **"Show me the code" mentality**

© 2018 National University of Singapore. All Rights Reserved

Fin